

# **JANUS FORWARD**

## **Wake State: Logic Trading Algorithm**

*Real-Time Decision Making, Pattern Recognition, and  
Trade Execution*

**Classification: Technical Implementation Guide**

**Version: 1.0 (Implementation-Ready)**

**Author:** Jordan Smith  
*github.com/nuniesmith*

**Date:** December 25, 2025

**JANUS Forward Overview:**

- **Purpose:** Real-time trading decisions during market hours
- **Hot Path:** Low-latency, high-throughput execution
- **Components:** Visual pattern recognition, symbolic reasoning, multimodal fusion, execution control
- **Goal:** Neuro-symbolic trading that combines deep learning with logical constraints

## Abstract

JANUS Forward represents the "wake state" of the JANUS trading system, responsible for all real-time decision-making during market hours. This document provides a comprehensive mathematical and implementation specification for the Forward service, which combines:

- **Visual Pattern Recognition** using Gramian Angular Fields (GAF) and Video Vision Transformers (ViViT)
- **Symbolic Reasoning** via Logic Tensor Networks (LTN) for constraint satisfaction
- **Multimodal Fusion** integrating time series, visual, and textual market data
- **Dual-Pathway Decision Making** inspired by basal ganglia architecture

The Forward service operates on a hot path with strict latency requirements, implementing end-to-end gradient flow through differentiable market simulation while maintaining regulatory compliance through symbolic constraints.

## Contents

<b>1 Visual Pattern Recognition: DiffGAF and ViViT</b>	<b>3</b>
1.1 Mathematical Foundation: Gramian Angular Fields . . . . .	3
1.1.1 Input Preprocessing . . . . .	3
1.1.2 Step 1: Learnable Normalization . . . . .	3
1.1.3 Step 2: Polar Coordinate Transformation . . . . .	3
1.1.4 Step 3: Gramian Field Generation . . . . .	3
1.1.5 Implementation Algorithm . . . . .	4
1.2 3D Spatiotemporal Manifolds: GAF Video . . . . .	4
1.2.1 Sliding Window GAF Video Generation . . . . .	4
1.2.2 Mathematical Formulation . . . . .	5
1.3 Video Vision Transformer (ViViT) . . . . .	5
1.3.1 Architecture Overview . . . . .	5
1.3.2 Patch Embedding . . . . .	5
1.3.3 Spatial Attention . . . . .	6
1.3.4 Temporal Attention . . . . .	6
1.3.5 Output Embedding . . . . .	6
<b>2 Logic Tensor Networks: Symbolic Reasoning Engine</b>	<b>6</b>
2.1 Mathematical Foundation . . . . .	6
2.1.1 Grounding Function . . . . .	6
2.1.2 Predicate Grounding . . . . .	6
2.2 Lukasiewicz T-Norm Operations . . . . .	7
2.2.1 Conjunction (AND) . . . . .	7
2.2.2 Disjunction (OR) . . . . .	7
2.2.3 Negation (NOT) . . . . .	7
2.2.4 Implication (IF-THEN) . . . . .	7
2.2.5 Universal Quantification (FOR ALL) . . . . .	7
2.2.6 Existential Quantification (EXISTS) . . . . .	7
2.3 Knowledge Base Formulation . . . . .	7
2.3.1 Wash Sale Constraint . . . . .	7
2.3.2 Almgren-Chriss Risk Constraint . . . . .	8
2.3.3 VPIN Toxicity Constraint . . . . .	8
2.4 Logical Loss Function . . . . .	8
2.4.1 Satisfiability Aggregation . . . . .	8
2.4.2 Logical Loss . . . . .	8
2.4.3 Combined Loss . . . . .	8

<b>3 Multimodal Fusion: Gated Cross-Attention</b>	<b>9</b>
3.1 Input Modalities . . . . .	9
3.2 Gated Cross-Attention Mechanism . . . . .	9
3.2.1 Attention Computation . . . . .	9
3.2.2 Gating Mechanism . . . . .	9
3.2.3 Fused Representation . . . . .	9
3.3 Multi-Modal Fusion Pipeline . . . . .	10
<b>4 Decision Engine: Basal Ganglia Pathways</b>	<b>10</b>
4.1 Praxeological Motor: Dual Pathways . . . . .	10
4.1.1 Direct Pathway (Go Signal) . . . . .	10
4.1.2 Indirect Pathway (No-Go Signal) . . . . .	10
4.1.3 Final Action . . . . .	10
4.2 Cerebellar Forward Model . . . . .	10
4.2.1 Market Impact Prediction . . . . .	11
4.2.2 Sensory Prediction Error . . . . .	11
4.2.3 Trajectory Adjustment . . . . .	11
<b>5 Implementation Checklist</b>	<b>12</b>
5.1 Core Components . . . . .	12
5.2 Integration & Testing . . . . .	13
5.3 Deployment Readiness . . . . .	14
<b>6 Rust Implementation Considerations</b>	<b>14</b>
6.1 Hot Path Optimization . . . . .	14
6.2 ML Framework Integration . . . . .	15
6.2.1 Option 1: PyTorch via tch-rs . . . . .	15
6.2.2 Option 2: ONNX Runtime via ort . . . . .	15
6.2.3 Option 3: Candle (Hugging Face) . . . . .	15
6.3 Error Handling Strategy . . . . .	15

# 1 Visual Pattern Recognition: DiffGAF and ViViT

The visual subsystem transforms time series data into spatiotemporal images, enabling the system to "see" market patterns that traditional numerical methods miss.

## 1.1 Mathematical Foundation: Gramian Angular Fields

### 1.1.1 Input Preprocessing

Given a univariate time series of length  $N$ :

$$X = \{x_1, x_2, \dots, x_N\} \in \mathbb{R}^N \quad (1)$$

### 1.1.2 Step 1: Learnable Normalization

The time series is normalized using learnable parameters  $\alpha \in \mathbb{R}^+$  and  $\beta \in \mathbb{R}$ :

$$\tilde{x}_i = \tanh \left( \frac{x_i - \min(X)}{\max(X) - \min(X) + \epsilon} \cdot \alpha + \beta \right) \quad (2)$$

where  $\epsilon = 10^{-8}$  prevents division by zero. The normalized values  $\tilde{x}_i \in [-1, 1]$  are constrained to the domain of the cosine function.

**Implementation Note:**  $\alpha$  and  $\beta$  are learnable parameters initialized as:

$$\alpha_0 = 1.0 \quad (3)$$

$$\beta_0 = 0.0 \quad (4)$$

These are optimized via backpropagation through the entire pipeline.

### 1.1.3 Step 2: Polar Coordinate Transformation

Each normalized value is mapped to polar coordinates:

$$\begin{cases} \phi_i = \arccos(\tilde{x}_i), & \tilde{x}_i \in [-1, 1] \\ r_i = \frac{i}{N}, & i \in \{1, 2, \dots, N\} \end{cases} \quad (5)$$

where  $\phi_i \in [0, \pi]$  is the angular component and  $r_i \in [0, 1]$  is the radial component (normalized time index).

### 1.1.4 Step 3: Gramian Field Generation

Two Gramian Angular Fields are computed:

**Gramian Angular Summation Field (GASF):**

$$\text{GASF}_{i,j} = \cos(\phi_i + \phi_j) = \tilde{x}_i \tilde{x}_j - \sqrt{1 - \tilde{x}_i^2} \sqrt{1 - \tilde{x}_j^2} \quad (6)$$

**Gramian Angular Difference Field (GADF):**

$$\text{GADF}_{i,j} = \sin(\phi_i - \phi_j) = \sqrt{1 - \tilde{x}_i^2} \tilde{x}_j - \tilde{x}_i \sqrt{1 - \tilde{x}_j^2} \quad (7)$$

The result is two  $N \times N$  matrices (images) where:

- The main diagonal ( $i = j$ ) contains the original normalized values
- Off-diagonal elements encode temporal correlations
- GASF captures summation relationships
- GADF captures difference relationships

**1.1.5 Implementation Algorithm****Algorithm 1** DiffGAF Transformation

**Require:** Time series  $X \in \mathbb{R}^N$ , learnable params  $\alpha, \beta$

**Ensure:** GASF and GADF matrices  $\in \mathbb{R}^{N \times N}$

```

1:  $X_{\min} \leftarrow \min(X)$ ,  $X_{\max} \leftarrow \max(X)$ 
2:  $\tilde{X} \leftarrow \tanh\left(\frac{X - X_{\min}}{X_{\max} - X_{\min} + \epsilon} \cdot \alpha + \beta\right)$ 
3:  $\Phi \leftarrow \arccos(\tilde{X})$  ▷ Element-wise arccos
4: Initialize GASF  $\leftarrow \mathbf{0}_{N \times N}$ , GADF  $\leftarrow \mathbf{0}_{N \times N}$ 
5: for  $i = 1$  to  $N$  do
6:   for  $j = 1$  to  $N$  do
7:      $\text{GASF}_{i,j} \leftarrow \cos(\Phi_i + \Phi_j)$ 
8:      $\text{GADF}_{i,j} \leftarrow \sin(\Phi_i - \Phi_j)$ 
9:   end for
10: end for
11: return GASF, GADF

```

**1.2 3D Spatiotemporal Manifolds: GAF Video****1.2.1 Sliding Window GAF Video Generation**

Given a time series  $X = \{x_1, x_2, \dots, x_T\}$  of length  $T$ , we generate a sequence of overlapping GAF frames:

$$\mathcal{V} = \{GAF(X_{t:t+w}), GAF(X_{t+s:t+w+s}), \dots, GAF(X_{t+(F-1)s:t+w+(F-1)s})\} \quad (8)$$

where:

- $w$  = window size (e.g., 60 timesteps)
- $s$  = stride (e.g., 10 timesteps)
- $F$  = number of frames (e.g., 16 frames)

Each frame is a  $2 \times N \times N$  tensor (GASF + GADF channels). The complete video tensor is:

$$\mathcal{V} \in \mathbb{R}^{F \times 2 \times N \times N} \quad (9)$$

## 1.2.2 Mathematical Formulation

For frame  $f \in \{0, 1, \dots, F - 1\}$ :

$$\mathcal{V}_f = \begin{bmatrix} \text{GASF}(X_{t+fs:t+w+fs}) \\ \text{GADF}(X_{t+fs:t+w+fs}) \end{bmatrix} \quad (10)$$

## 1.3 Video Vision Transformer (ViViT)

### 1.3.1 Architecture Overview

ViViT processes the 3D tensor  $\mathcal{V}$  using factorized spatial-temporal attention.

### 1.3.2 Patch Embedding

Each frame is divided into patches. For a frame of size  $H \times W$  with patch size  $P$ :

$$N_p = \frac{H}{P} \times \frac{W}{P} \quad (11)$$

patches per frame.

The patch embedding for patch  $(i, j)$  in frame  $f$  is:

$$\mathbf{z}_{f,i,j}^{(0)} = \mathbf{E} \cdot \text{flatten}(\mathcal{V}_{f,i:i+P,j:j+P}) + \mathbf{p}_{f,i,j} \quad (12)$$

where:

- $\mathbf{E} \in \mathbb{R}^{d \times (2P^2)}$  is the embedding matrix
- $\mathbf{p}_{f,i,j}$  is the positional encoding (spatial + temporal)

### 1.3.3 Spatial Attention

Within each frame  $f$ , spatial self-attention is computed:

$$\text{Attention}(\mathbf{Q}_s, \mathbf{K}_s, \mathbf{V}_s) = \text{softmax}\left(\frac{\mathbf{Q}_s \mathbf{K}_s^\top}{\sqrt{d_h}}\right) \mathbf{V}_s \quad (13)$$

where  $\mathbf{Q}_s, \mathbf{K}_s, \mathbf{V}_s$  are queries, keys, and values from spatial patches.

### 1.3.4 Temporal Attention

Across frames, temporal attention captures evolution:

$$\text{Attention}(\mathbf{Q}_t, \mathbf{K}_t, \mathbf{V}_t) = \text{softmax}\left(\frac{\mathbf{Q}_t \mathbf{K}_t^\top}{\sqrt{d_h}}\right) \mathbf{V}_t \quad (14)$$

### 1.3.5 Output Embedding

The final visual embedding vector is:

$$\mathbf{e}_{\text{visual}} = \text{MLP}(\text{GlobalPool}(\mathbf{Z}^{(L)})) \in \mathbb{R}^{d_{\text{embed}}} \quad (15)$$

where  $L$  is the number of transformer layers and  $d_{\text{embed}}$  is the embedding dimension (e.g., 768).

## 2 Logic Tensor Networks: Symbolic Reasoning Engine

The LTN subsystem ensures that all trading decisions satisfy regulatory and risk management constraints through differentiable first-order logic.

### 2.1 Mathematical Foundation

#### 2.1.1 Grounding Function

Let  $\mathcal{G} : \mathcal{S} \rightarrow \mathbb{R}^n$  be a grounding function that maps logical symbols to tensors:

$$\mathcal{G} : \text{Constants} \cup \text{Predicates} \cup \text{Functions} \rightarrow \mathbb{R}^n \quad (16)$$

#### 2.1.2 Predicate Grounding

A predicate  $P$  with arity  $k$  is grounded as a neural network:

$$\mathcal{G}(P) : \mathbb{R}^{k \times d} \rightarrow [0, 1] \quad (17)$$

where  $d$  is the embedding dimension.

For example,  $\text{IsVolatile}(\text{market})$  is implemented as:

$$\mathcal{G}(\text{IsVolatile})(\mathbf{e}_{\text{market}}) = \text{sigmoid}(\mathbf{W}_v \mathbf{e}_{\text{market}} + b_v) \quad (18)$$

where  $\mathbf{W}_v \in \mathbb{R}^{1 \times d}$  and  $b_v \in \mathbb{R}$  are learnable parameters.

## 2.2 Lukasiewicz T-Norm Operations

### 2.2.1 Conjunction (AND)

$$\mathcal{G}(A \wedge B) = \max(0, \mathcal{G}(A) + \mathcal{G}(B) - 1) \quad (19)$$

### 2.2.2 Disjunction (OR)

$$\mathcal{G}(A \vee B) = \min(1, \mathcal{G}(A) + \mathcal{G}(B)) \quad (20)$$

### 2.2.3 Negation (NOT)

$$\mathcal{G}(\neg A) = 1 - \mathcal{G}(A) \quad (21)$$

### 2.2.4 Implication (IF-THEN)

$$\mathcal{G}(A \rightarrow B) = \min(1, 1 - \mathcal{G}(A) + \mathcal{G}(B)) \quad (22)$$

### 2.2.5 Universal Quantification (FOR ALL)

For a formula  $\phi(x)$  with free variable  $x$ :

$$\mathcal{G}(\forall x : \phi(x)) = \min_{x \in \mathcal{D}} \mathcal{G}(\phi(x)) \quad (23)$$

where  $\mathcal{D}$  is the domain of  $x$ .

### 2.2.6 Existential Quantification (EXISTS)

$$\mathcal{G}(\exists x : \phi(x)) = \max_{x \in \mathcal{D}} \mathcal{G}(\phi(x)) \quad (24)$$

## 2.3 Knowledge Base Formulation

### 2.3.1 Wash Sale Constraint

The wash sale rule is encoded as:

$$\forall t, \forall k \in [1, 30] : \neg(\text{SaleAtLoss}(t) \wedge \text{Buy}(t + k)) \quad (25)$$

In grounded form:

$$\mathcal{G}(\text{WashSale}) = \min_{t,k \in [1,30]} [1 - \max(0, \mathcal{G}(\text{SaleAtLoss})(t) + \mathcal{G}(\text{Buy})(t+k) - 1)] \quad (26)$$

### 2.3.2 Almgren-Chriss Risk Constraint

$$\forall v : \text{Volatile}(Market) \rightarrow \text{Impact}(v) < \text{Threshold}(\sigma) \quad (27)$$

Grounded:

$$\mathcal{G}(\text{ACRisk}) = \min_v [\min(1, 1 - \mathcal{G}(\text{Volatile}) + \mathcal{G}(\text{Impact}(v) < \text{Threshold}(\sigma)))] \quad (28)$$

where:

$$\text{Threshold}(\sigma) = \eta \cdot \sigma \cdot \sqrt{\frac{v}{V}} \quad (29)$$

with  $\eta$  = impact coefficient,  $\sigma$  = volatility,  $v$  = trade size,  $V$  = average volume.

### 2.3.3 VPIN Toxicity Constraint

$$\forall t : \text{VPIN}_t > \tau_{\text{VPIN}} \rightarrow \text{HaltTrading}(t) \quad (30)$$

Grounded:

$$\mathcal{G}(\text{VPIN}) = \min_t [\min(1, 1 - \mathcal{G}(\text{VPIN}_t > \tau_{\text{VPIN}}) + \mathcal{G}(\text{HaltTrading})(t))] \quad (31)$$

## 2.4 Logical Loss Function

### 2.4.1 Satisfiability Aggregation

For a knowledge base  $\mathcal{K} = \{\phi_1, \phi_2, \dots, \phi_m\}$ :

$$\text{SatAgg}(\mathcal{K}) = \left( \frac{1}{m} \sum_{i=1}^m \mathcal{G}(\phi_i)^p \right)^{1/p} \quad (32)$$

where  $p$  is the generalized mean parameter (typically  $p = 2$  for quadratic mean).

### 2.4.2 Logical Loss

$$\mathcal{L}_{\text{logic}}(\theta) = 1 - \text{SatAgg}(\mathcal{K}) \quad (33)$$

### 2.4.3 Combined Loss

The total loss combines predictive and logical components:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{predictive}} + \lambda_{\text{logic}} \cdot \mathcal{L}_{\text{logic}} \quad (34)$$

where  $\lambda_{\text{logic}}$  is a hyperparameter (typically 0.1 to 1.0).

## 3 Multimodal Fusion: Gated Cross-Attention

The fusion subsystem integrates visual, temporal, and textual market information into a unified representation.

### 3.1 Input Modalities

The system receives three input streams:

$$\mathbf{H}_{\text{TS}} \in \mathbb{R}^{L_{\text{TS}} \times d} \quad (\text{Time Series Tokens from Chronos-Bolt}) \quad (35)$$

$$\mathbf{H}_{\text{Vis}} \in \mathbb{R}^{L_{\text{Vis}} \times d} \quad (\text{Visual Embeddings from ViViT}) \quad (36)$$

$$\mathbf{H}_{\text{Text}} \in \mathbb{R}^{L_{\text{Text}} \times d} \quad (\text{Text Embeddings from FinBERT}) \quad (37)$$

### 3.2 Gated Cross-Attention Mechanism

#### 3.2.1 Attention Computation

For primary modality  $m$  and auxiliary modality  $n$ :

$$\alpha_{m \rightarrow n} = \text{softmax} \left( \frac{\mathbf{Q}_m \mathbf{K}_n^\top}{\sqrt{d_h}} \right) \quad (38)$$

where:

$$\mathbf{Q}_m = \mathbf{H}_m \mathbf{W}_Q \quad (39)$$

$$\mathbf{K}_n = \mathbf{H}_n \mathbf{W}_K \quad (40)$$

$$\mathbf{V}_n = \mathbf{H}_n \mathbf{W}_V \quad (41)$$

#### 3.2.2 Gating Mechanism

The gating scalar is computed as:

$$\lambda_{\text{gate}} = \text{sigmoid}(\mathbf{W}_g[\mathbf{H}_m; \mathbf{H}_n] + b_g) \quad (42)$$

where  $[\cdot; \cdot]$  denotes concatenation.

#### 3.2.3 Fused Representation

$$\mathbf{H}_{\text{fused}} = \mathbf{H}_m + \lambda_{\text{gate}} \cdot \alpha_{m \rightarrow n} \mathbf{V}_n \quad (43)$$

### 3.3 Multi-Modal Fusion Pipeline

The complete fusion process:

$$\mathbf{H}_1 = \mathbf{H}_{\text{TS}} + \lambda_{\text{vis}} \cdot \text{CrossAttn}(\mathbf{H}_{\text{TS}}, \mathbf{H}_{\text{Vis}}) \quad (44)$$

$$\mathbf{H}_2 = \mathbf{H}_1 + \lambda_{\text{text}} \cdot \text{CrossAttn}(\mathbf{H}_1, \mathbf{H}_{\text{Text}}) \quad (45)$$

$$\mathbf{e}_{\text{final}} = \text{GlobalPool}(\mathbf{H}_2) \quad (46)$$

## 4 Decision Engine: Basal Ganglia Pathways

The decision engine implements a dual-pathway architecture inspired by the basal ganglia, with separate "go" and "no-go" pathways.

### 4.1 Praxeological Motor: Dual Pathways

#### 4.1.1 Direct Pathway (Go Signal)

The direct pathway generates action proposals:

$$a_{\text{proposed}} = \arg \max_{a \in \mathcal{A}} [\mathbf{W}_{\alpha} \mathbf{e}_{\text{final}} + b_{\alpha}]_a \quad (47)$$

where  $\mathcal{A}$  is the action space (BUY, SELL, HOLD, or continuous trade sizes).

#### 4.1.2 Indirect Pathway (No-Go Signal)

The indirect pathway computes risk veto:

$$v_{\text{risk}} = \text{sigmoid}(\mathbf{W}_{\text{risk}}[\mathbf{e}_{\text{final}}; \text{VPIN}; \sigma_{\text{market}}] + b_{\text{risk}}) \quad (48)$$

#### 4.1.3 Final Action

The final action is gated by the risk veto:

$$a_{\text{final}} = \begin{cases} a_{\text{proposed}} & \text{if } v_{\text{risk}} < \tau_{\text{risk}} \text{ AND } \mathcal{L}_{\text{logic}} < \tau_{\text{logic}} \\ \text{HOLD} & \text{otherwise} \end{cases} \quad (49)$$

where  $\tau_{\text{risk}}$  is the risk threshold (e.g., 0.7) and  $\tau_{\text{logic}}$  is the logical constraint threshold (e.g., 0.1).

### 4.2 Cerebellar Forward Model

#### 4.2.1 Market Impact Prediction

The forward model predicts execution price:

$$\hat{p}_{\text{exec}} = f_{\text{forward}}(\mathbf{s}_{\text{LOB}}, v, a_{\text{final}}) \quad (50)$$

where  $\mathbf{s}_{\text{LOB}}$  is the limit order book state.

#### 4.2.2 Sensory Prediction Error

$$\text{SPE} = |p_{\text{actual}} - \hat{p}_{\text{exec}}| \quad (51)$$

#### 4.2.3 Trajectory Adjustment

The execution trajectory is adjusted:

$$v_{\text{adjusted}} = v_{\text{original}} - \eta_{\text{cerebellar}} \cdot \text{SPE} \cdot \nabla_v \text{SPE} \quad (52)$$

## 5 Implementation Checklist

sec:checklist

This section provides a sequential checklist for implementing JANUS Forward.

### 5.1 Core Components

#### 1. Visual Pattern Recognition Module

- Implement DiffGAF normalization with learnable  $\alpha, \beta$
- Implement polar coordinate transformation
- Implement GASF and GADF computation
- Implement sliding window GAF video generation
- Implement ViViT patch embedding
- Implement spatial attention mechanism
- Implement temporal attention mechanism
- Test on sample time series data
- Benchmark latency (target: <50ms for inference)

#### 2. Logic Tensor Networks Module

- Implement grounding function framework
- Implement predicate neural networks
- Implement Lukasiewicz T-norm operations
- Implement universal/existential quantification
- Encode wash sale constraint
- Encode Almgren-Chriss constraint
- Encode VPIN constraint
- Implement satisfiability aggregation
- Implement logical loss function
- Test constraint satisfaction with edge cases
- Benchmark latency (target: <10ms for evaluation)

#### 3. Multimodal Fusion Module

- Implement time series tokenization (Chronos-Bolt integration)
- Implement visual embedding extraction

- Implement text embedding (FinBERT integration)
- Implement gated cross-attention mechanism
- Implement multi-modal fusion pipeline
- Test fusion on sample multimodal data
- Validate attention weight distributions

#### 4. Decision Engine Module

- Implement direct pathway (alpha motor)
- Implement indirect pathway (risk motor)
- Implement risk-gated action selection
- Implement logic-gated action selection
- Implement cerebellar forward model
- Implement sensory prediction error computation
- Implement trajectory adjustment
- Test end-to-end decision making
- Validate constraint adherence in production scenarios

## 5.2 Integration & Testing

### 1. End-to-End Pipeline

- Connect all modules into unified forward pass
- Implement gradient flow verification
- Test backpropagation through entire pipeline
- Validate differentiable constraint satisfaction

### 2. Performance Optimization

- Profile latency bottlenecks
- Optimize tensor operations for GPU
- Implement model quantization (INT8/FP16)
- Add batching support for parallel inference
- Target: <100ms end-to-end latency

### 3. Safety & Validation

- Add input validation and sanitization

- Implement kill switch integration
- Add logging for all trading decisions
- Implement emergency halt on constraint violation
- Test failure modes (network outage, invalid data, etc.)

## 5.3 Deployment Readiness

### 1. Production Hardening

- Remove all panic!() calls
- Replace unwrap() with proper error handling
- Add comprehensive error types
- Implement graceful degradation
- Add health check endpoints

### 2. Monitoring & Observability

- Add metrics export (Prometheus format)
- Implement distributed tracing
- Log all constraint violations
- Monitor inference latency
- Track prediction accuracy

## 6 Rust Implementation Considerations

sec:rust

### 6.1 Hot Path Optimization

The Forward service must maintain low latency (<100ms end-to-end). Key Rust optimizations:

- **Zero-copy operations:** Use ndarray views instead of clones
- **SIMD acceleration:** Leverage packed\_simd for GAF computation
- **Async runtime:** Use tokio for non-blocking I/O
- **Memory pooling:** Pre-allocate tensor buffers to avoid allocation overhead

## 6.2 ML Framework Integration

### 6.2.1 Option 1: PyTorch via tch-rs

- Pros: Full PyTorch ecosystem, easy model export
- Cons: Requires LibTorch, larger binary size

### 6.2.2 Option 2: ONNX Runtime via ort

- Pros: Lightweight, cross-platform, optimized inference
- Cons: Limited to inference, requires model conversion
- **Recommended for production**

### 6.2.3 Option 3: Candle (Hugging Face)

- Pros: Pure Rust, no C++ dependencies
- Cons: Younger ecosystem, fewer pre-trained models
- **Recommended for future migration**

## 6.3 Error Handling Strategy

```

1 // Custom error types for Forward service
2 #[derive(Debug, thiserror::Error)]
3 pub enum ForwardError {
4     #[error("GAF transformation failed: {0}")]
5     GafError(String),
6
7     #[error("LTN constraint violation: {constraint}")]
8     ConstraintViolation { constraint: String },
9
10    #[error("Model inference failed: {0}")]
11    InferenceError(String),
12
13    #[error("Risk threshold exceeded: {risk_score}")]
14    RiskVeto { risk_score: f64 },
15 }
16
17 // Result type alias
18 pub type ForwardResult<T> = Result<T, ForwardError>;

```

## References

- [1] Jordan Smith, "Project JANUS: Implementation Guide v1.0," 2025.
- [2] Wang, Oates, "Encoding Time Series as Images for Visual Inspection and Classification Using Tiled Convolutional Neural Networks," AAAI 2015.
- [3] Arnab et al., "ViViT: A Video Vision Transformer," ICCV 2021.
- [4] Badreddine et al., "Logic Tensor Networks," Artificial Intelligence, 2022.
- [5] Ansari et al., "Chronos: Learning the Language of Time Series," arXiv:2403.07815, 2024.
- [6] Araci, "FinBERT: Financial Sentiment Analysis with Pre-trained Language Models," arXiv:1908.10063, 2019.
- [7] Almgren, Chriss, "Optimal Execution of Portfolio Transactions," Journal of Risk, 2000.
- [8] Easley et al., "Flow Toxicity and Liquidity in a High-frequency World," Review of Financial Studies, 2012.