

# Project JANUS

Neuromorphic Trading Intelligence

## Complete Technical Specification

*A Brain-Inspired Architecture for Autonomous Financial Systems*

### Unified Documentation

This document consolidates all technical specifications of Project JANUS:

1. **Main Architecture** — System design and philosophical foundation
2. **Forward Service** — Real-time decision-making and execution
3. **Backward Service** — Memory consolidation and learning
4. **Neuromorphic Architecture** — Brain-region mapping
5. **Rust Implementation** — Production deployment guide

**Author**

Jordan Smith

**Date**

December 28, 2025

*"The god of beginnings and transitions, looking simultaneously to the future and the past."*

## Contents

<b>I Main Architecture</b>	<b>2</b>
<b>II Forward Service (Janus Bifrons)</b>	<b>3</b>
<b>1 Visual Pattern Recognition: DiffGAF and ViViT</b>	<b>3</b>
1.1 Mathematical Foundation: Gramian Angular Fields . . . . .	3
1.1.1 Input Preprocessing . . . . .	3
1.1.2 Step 1: Learnable Normalization . . . . .	3
1.1.3 Step 2: Polar Coordinate Transformation . . . . .	4
1.1.4 Step 3: Gramian Field Generation . . . . .	4
1.2 3D Spatiotemporal Manifolds: GAF Video . . . . .	4
1.2.1 Sliding Window GAF Video Generation . . . . .	4
1.3 Video Vision Transformer (ViViT) . . . . .	4
1.3.1 Patch Embedding . . . . .	5
1.3.2 Spatial Attention . . . . .	5
1.3.3 Temporal Attention . . . . .	5
<b>2 Logic Tensor Networks: Symbolic Reasoning Engine</b>	<b>5</b>
2.1 Mathematical Foundation . . . . .	5
2.1.1 Grounding Function . . . . .	5
2.1.2 Predicate Grounding . . . . .	5
2.2 Lukasiewicz T-Norm Operations . . . . .	5
2.2.1 Conjunction (AND) . . . . .	5
2.2.2 Disjunction (OR) . . . . .	6
2.2.3 Negation (NOT) . . . . .	6
2.2.4 Implication (IF-THEN) . . . . .	6
2.3 Knowledge Base Formulation . . . . .	6
2.3.1 Wash Sale Constraint . . . . .	6
2.3.2 Almgren-Chriss Risk Constraint . . . . .	6
2.4 Logical Loss Function . . . . .	6
2.4.1 Satisfiability Aggregation . . . . .	6
2.4.2 Logical Loss . . . . .	6
<b>3 Multimodal Fusion: Gated Cross-Attention</b>	<b>6</b>
3.1 Input Modalities . . . . .	7
3.2 Gated Cross-Attention Mechanism . . . . .	7

3.2.1	Attention Computation . . . . .	7
3.2.2	Gating Mechanism . . . . .	7
<b>4</b>	<b>Decision Engine: Basal Ganglia Pathways</b>	<b>7</b>
4.1	Praxeological Motor: Dual Pathways . . . . .	7
4.1.1	Direct Pathway (Go Signal) . . . . .	7
4.1.2	Indirect Pathway (No-Go Signal) . . . . .	7
4.2	Cerebellar Forward Model . . . . .	7
4.2.1	Market Impact Prediction . . . . .	7
<b>III</b>	<b>Backward Service (Janus Consivius)</b>	<b>8</b>
<b>5</b>	<b>Memory Hierarchy: Three-Timescale Architecture</b>	<b>8</b>
5.1	Short-Term Memory (Hippocampus) . . . . .	8
5.1.1	Episodic Buffer . . . . .	8
5.1.2	Pattern Separation . . . . .	8
5.2	Medium-Term Consolidation (SWR Simulator) . . . . .	9
5.2.1	Replay Prioritization . . . . .	9
5.2.2	Sampling Probability . . . . .	9
5.2.3	Importance Sampling Correction . . . . .	9
5.3	Long-Term Memory (Neocortex) . . . . .	9
5.3.1	Schema Representation . . . . .	9
5.3.2	Recall-Gated Consolidation . . . . .	9
<b>6</b>	<b>UMAP Visualization: Cognitive Dashboard</b>	<b>9</b>
6.1	AlignedUMAP for Schema Formation . . . . .	9
6.1.1	Objective Function . . . . .	10
6.2	Parametric UMAP for Real-Time Monitoring . . . . .	10
<b>7</b>	<b>Integration with Vector Database (Qdrant)</b>	<b>10</b>
7.1	Schema Storage . . . . .	10
7.2	Similarity Search . . . . .	11
<b>IV</b>	<b>Neuromorphic Architecture</b>	<b>12</b>
<b>8</b>	<b>Neuromorphic Design Philosophy</b>	<b>12</b>
8.1	Why Brain-Inspired Architecture? . . . . .	12
8.2	Neuroscience-to-Trading Mapping . . . . .	12

<b>9 Brain Region Architectures</b>	<b>13</b>
9.1 Cortex: Strategic Planning & Long-term Memory . . . . .	13
9.1.1 Trading Implementation . . . . .	13
9.2 Hippocampus: Episodic Memory & Experience Replay . . . . .	13
9.2.1 Trading Implementation . . . . .	13
9.3 Basal Ganglia: Action Selection & Reinforcement Learning . . . . .	13
9.3.1 Trading Implementation . . . . .	13
9.4 Prefrontal Cortex: Logic, Planning & Compliance . . . . .	14
9.4.1 Trading Implementation . . . . .	14
9.5 Amygdala: Fear, Threat Detection & Circuit Breakers . . . . .	14
9.5.1 Trading Implementation . . . . .	14
9.6 Cerebellum: Motor Control & Execution . . . . .	14
9.6.1 Trading Implementation . . . . .	14
<b>V Rust Implementation</b>	<b>16</b>
<b>10 Architectural Overview</b>	<b>16</b>
10.1 The Rust-First Philosophy . . . . .	16
10.2 Component Diagram . . . . .	17
<b>11 Machine Learning Framework Strategy</b>	<b>17</b>
11.1 Framework Comparison Matrix . . . . .	17
11.2 Recommended Migration Path . . . . .	17
11.2.1 Phase 1: Hybrid (Months 1-3) . . . . .	17
11.2.2 Phase 2: Rust-Native Inference (Months 4-6) . . . . .	18
11.2.3 Phase 3: Full Rust ML (Months 7-12) . . . . .	18
<b>12 Forward Service: Rust Implementation</b>	<b>18</b>
12.1 Performance Requirements . . . . .	18
12.2 Core Data Structures . . . . .	18
12.3 GAF Transformation Algorithm . . . . .	19
12.4 LTN Constraint Evaluation . . . . .	19
12.5 Async Service Architecture . . . . .	20
<b>13 Backward Service: Batch Processing</b>	<b>20</b>
13.1 Prioritized Experience Replay . . . . .	20
13.2 Schema Consolidation Algorithm . . . . .	21
<b>14 Deployment Architecture</b>	<b>22</b>
14.1 Service Orchestration . . . . .	22

## Part I

# Main Architecture

## Overview

Part 1 provides the architectural philosophy and system design overview for Project JANUS. This section would typically include:

- **Introduction:** The crisis of complexity in quantitative trading
- **Architectural Philosophy:** The dual-service design (Forward/Backward)
- **Core Components:** Vision, Logic, Fusion, and Decision systems
- **Memory Hierarchy:** Three-timescale learning architecture
- **Implementation Strategy:** Rust-first development approach

**Note:** The detailed mathematical specifications for each component are presented in Parts 2-5 below. This consolidated document focuses on the theoretical foundations and algorithmic specifications rather than high-level architectural discussion.

## Part II

# Forward Service (Janus Bifrons)

## Abstract

JANUS Forward represents the "wake state" of the JANUS trading system, responsible for all real-time decision-making during market hours. This service combines:

- **Visual Pattern Recognition** using Gramian Angular Fields (GAF) and Video Vision Transformers (ViViT)
- **Symbolic Reasoning** via Logic Tensor Networks (LTN) for constraint satisfaction
- **Multimodal Fusion** integrating time series, visual, and textual market data
- **Dual-Pathway Decision Making** inspired by basal ganglia architecture

The Forward service operates on a hot path with strict latency requirements, implementing end-to-end gradient flow through differentiable market simulation while maintaining regulatory compliance through symbolic constraints.

## 1 Visual Pattern Recognition: DiffGAF and ViViT

The visual subsystem transforms time series data into spatiotemporal images, enabling the system to "see" market patterns that traditional numerical methods miss.

### 1.1 Mathematical Foundation: Gramian Angular Fields

Time series are encoded into polar coordinates and projected onto Gramian matrices, creating 2D representations that preserve temporal correlations.

#### 1.1.1 Input Preprocessing

Given raw market data  $X = \{x_1, x_2, \dots, x_T\}$  where  $x_t \in \mathbb{R}^D$  (multi-feature time series), we first apply feature selection to extract  $F$  relevant features.

#### 1.1.2 Step 1: Learnable Normalization

Instead of fixed min-max scaling, we use learnable affine transformations with domain constraints:

$$\tilde{x}_t = \tanh \left( \gamma \odot \frac{x_t - \mu}{\sigma} + \beta \right) \quad (1)$$

where  $\gamma, \beta \in \mathbb{R}^F$  are learned parameters, and  $\mu, \sigma$  are running statistics. The  $\tanh$  function guarantees  $\tilde{x}_t \in (-1, 1)$ , ensuring the subsequent  $\arccos$  operation is well-defined.

### 1.1.3 Step 2: Polar Coordinate Transformation

Map normalized values to angular space:

$$\phi_t = \arccos(\tilde{x}_t) \in [0, \pi] \quad (2)$$

$$r_t = \frac{t}{T} \quad (\text{normalized timestamp}) \quad (3)$$

### 1.1.4 Step 3: Gramian Field Generation

Construct the Gramian Angular Summation Field (GASF):

$$\mathbf{G}_{ij} = \cos(\phi_i + \phi_j) = \tilde{x}_i \tilde{x}_j - \sqrt{1 - \tilde{x}_i^2} \sqrt{1 - \tilde{x}_j^2} \quad (4)$$

Or the Gramian Angular Difference Field (GADF):

$$\mathbf{G}_{ij} = \sin(\phi_i - \phi_j) = \sqrt{1 - \tilde{x}_i^2} \tilde{x}_j - \tilde{x}_i \sqrt{1 - \tilde{x}_j^2} \quad (5)$$

## 1.2 3D Spatiotemporal Manifolds: GAF Video

To capture temporal dynamics, we generate a sequence of GAF frames using sliding windows.

### 1.2.1 Sliding Window GAF Video Generation

Given a time series of length  $T$ , window size  $W$ , and stride  $S$ :

1. Extract windows:  $X_k = \{x_{(k-1)S+1}, \dots, x_{(k-1)S+W}\}$  for  $k = 1, \dots, N$
2. Generate GAF for each window:  $\mathbf{G}_k = \text{GAF}(X_k) \in \mathbb{R}^{W \times W}$
3. Stack into video:  $\mathbf{V} = [\mathbf{G}_1, \mathbf{G}_2, \dots, \mathbf{G}_N] \in \mathbb{R}^{N \times W \times W}$

## 1.3 Video Vision Transformer (ViViT)

The GAF video is processed by a factorized spatiotemporal transformer.

### 1.3.1 Patch Embedding

Divide each frame  $\mathbf{G}_k$  into non-overlapping patches:

$$\mathbf{P}_k = \text{Reshape}(\mathbf{G}_k) \in \mathbb{R}^{P \times (p^2)} \quad (6)$$

where  $P = (W/p)^2$  is the number of patches per frame.

### 1.3.2 Spatial Attention

Apply self-attention within each frame:

$$\mathbf{Z}_k^{(l)} = \text{MSA}(\text{LN}(\mathbf{Z}_k^{(l-1)})) + \mathbf{Z}_k^{(l-1)} \quad (7)$$

### 1.3.3 Temporal Attention

Apply attention across frames:

$$\mathbf{H}^{(l)} = \text{MSA}(\text{LN}([\mathbf{Z}_1^{(l)}, \dots, \mathbf{Z}_N^{(l)}])) \quad (8)$$

## 2 Logic Tensor Networks: Symbolic Reasoning Engine

LTNs bridge neural networks and first-order logic, enabling differentiable constraint satisfaction.

### 2.1 Mathematical Foundation

#### 2.1.1 Grounding Function

Map logical constants to real vectors:

$$\mathcal{G} : \mathcal{C} \rightarrow \mathbb{R}^d \quad (9)$$

#### 2.1.2 Predicate Grounding

A predicate  $P(x)$  is grounded as a neural network  $f_\theta : \mathbb{R}^d \rightarrow [0, 1]$ .

### 2.2 Lukasiewicz T-Norm Operations

#### 2.2.1 Conjunction (AND)

For training, we use Product Logic to ensure smooth gradients:

$$u \wedge v = u \cdot v \quad (10)$$

For inference/evaluation, standard Łukasiewicz logic may be used:

$$u \wedge v = \max(0, u + v - 1) \quad (11)$$

### 2.2.2 Disjunction (OR)

$$u \vee v = \min(1, u + v) \quad (12)$$

### 2.2.3 Negation (NOT)

$$\neg u = 1 - u \quad (13)$$

### 2.2.4 Implication (IF-THEN)

For training (Product Logic):

$$u \Rightarrow v = 1 - u + u \cdot v \quad (14)$$

For inference (Łukasiewicz Logic):

$$u \Rightarrow v = \min(1, 1 - u + v) \quad (15)$$

## 2.3 Knowledge Base Formulation

### 2.3.1 Wash Sale Constraint

$$\forall t : \text{Sell}(t) \wedge \text{Buy}(t') \wedge |t - t'| < 30 \Rightarrow \neg \text{TaxLoss}(t) \quad (16)$$

### 2.3.2 Almgren-Chriss Risk Constraint

$$\forall \text{order} : \text{Execute}(\text{order}) \Rightarrow \text{Slippage}(\text{order}) < \lambda \cdot \text{Volatility} \quad (17)$$

## 2.4 Logical Loss Function

### 2.4.1 Satisfiability Aggregation

$$\text{SAT}(\mathcal{KB}) = \text{p-mean}_{i=1}^{|\mathcal{KB}|}(\phi_i) \quad (18)$$

### 2.4.2 Logical Loss

$$\mathcal{L}_{\text{logic}} = 1 - \text{SAT}(\mathcal{KB}) \quad (19)$$

## 3 Multimodal Fusion: Gated Cross-Attention

### 3.1 Input Modalities

- Visual:  $\mathbf{v} \in \mathbb{R}^{d_v}$  (from ViViT)
- Temporal:  $\mathbf{t} \in \mathbb{R}^{d_t}$  (from LSTM/Transformer)
- Textual:  $\mathbf{s} \in \mathbb{R}^{d_s}$  (from BERT embeddings)

### 3.2 Gated Cross-Attention Mechanism

#### 3.2.1 Attention Computation

$$\text{Attn}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right)\mathbf{V} \quad (20)$$

#### 3.2.2 Gating Mechanism

$$g = \text{sigmoid}(\mathbf{W}_g[\mathbf{v}; \mathbf{t}; \mathbf{s}] + \mathbf{b}_g) \quad (21)$$

## 4 Decision Engine: Basal Ganglia Pathways

### 4.1 Praxeological Motor: Dual Pathways

#### 4.1.1 Direct Pathway (Go Signal)

$$\mathbf{d}_{\text{direct}} = \text{ReLU}(\mathbf{W}_d \mathbf{h}_{\text{fused}} + \mathbf{b}_d) \quad (22)$$

#### 4.1.2 Indirect Pathway (No-Go Signal)

$$\mathbf{d}_{\text{indirect}} = \text{ReLU}(\mathbf{W}_i \mathbf{h}_{\text{fused}} + \mathbf{b}_i) \quad (23)$$

### 4.2 Cerebellar Forward Model

#### 4.2.1 Market Impact Prediction

$$\hat{p}_{t+1} = f_{\text{cerebellum}}(\mathbf{s}_t, \mathbf{a}_t) \quad (24)$$

## Part III

# Backward Service (Janus Consivius)

## Abstract

JANUS Backward represents the "sleep state" of the system, responsible for memory consolidation, schema formation, and learning from accumulated experience. This service implements:

- **Three-Timescale Memory Hierarchy** (Hippocampus → SWR → Neocortex)
- **Sharp-Wave Ripple Simulation** for prioritized experience replay
- **Schema Formation** via UMAP-based clustering
- **Recall-Gated Consolidation** ensuring only successful patterns are promoted

The Backward service runs on a cold path during off-market hours, performing computationally intensive operations to distill daily experiences into long-term knowledge.

## 5 Memory Hierarchy: Three-Timescale Architecture

### 5.1 Short-Term Memory (Hippocampus)

#### 5.1.1 Episodic Buffer

Stores raw experiences during trading:

$$\mathcal{D}_{\text{hippo}} = \{(s_t, a_t, r_t, s_{t+1}, \mathbf{c}_t)\}_{t=1}^T \quad (25)$$

where  $\mathbf{c}_t$  contains contextual metadata (volatility, spreads, volume).

#### 5.1.2 Pattern Separation

Uses random projections to ensure diverse encoding:

$$\mathbf{h}_t = \tanh(\mathbf{W}_{\text{rand}} \cdot [s_t; a_t; \mathbf{c}_t]) \quad (26)$$

## 5.2 Medium-Term Consolidation (SWR Simulator)

### 5.2.1 Replay Prioritization

Compute TD-error based priority:

$$p_i = |\delta_i| + \epsilon \quad (27)$$

where  $\delta_i = r_i + \gamma \max_{a'} Q(s_{i+1}, a') - Q(s_i, a_i)$  and  $\epsilon = 10^{-6}$  ensures numerical stability.

### 5.2.2 Sampling Probability

$$P(i) = \frac{p_i^\alpha}{\sum_j p_j^\alpha} \quad (28)$$

where  $\alpha \in [0, 1]$  controls prioritization strength (typically  $\alpha = 0.6$ ).

### 5.2.3 Importance Sampling Correction

$$w_i = \left( \frac{1}{N \cdot P(i)} \right)^\beta \quad (29)$$

where  $\beta \in [0, 1]$  is annealed from  $0.4 \rightarrow 1.0$  during training to fully correct bias at convergence.

## 5.3 Long-Term Memory (Neocortex)

### 5.3.1 Schema Representation

Each schema is a prototype:

$$\mathbf{z}_k = \frac{1}{|\mathcal{C}_k|} \sum_{i \in \mathcal{C}_k} \mathbf{h}_i \quad (30)$$

### 5.3.2 Recall-Gated Consolidation

Only update schemas from successfully recalled experiences:

$$\mathbf{z}_k \leftarrow \mathbf{z}_k + \eta \cdot \mathbb{1}[\text{recall\_success}] \cdot (\mathbf{h}_{\text{new}} - \mathbf{z}_k) \quad (31)$$

## 6 UMAP Visualization: Cognitive Dashboard

### 6.1 AlignedUMAP for Schema Formation

Maintains consistent embeddings across sleep cycles.

### 6.1.1 Objective Function

The full UMAP loss includes both attraction and repulsion terms:

$$\mathcal{L}_{\text{UMAP}} = \sum_{i \neq j} [w_{ij} \log(q_{ij}) + (1 - w_{ij}) \log(1 - q_{ij})] \quad (32)$$

where  $q_{ij} = (1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}$ .

**Note:** In practice, the repulsion term  $(1 - w_{ij})$  is approximated via *negative sampling* to achieve  $\mathcal{O}(N)$  complexity. For each positive edge, we sample  $k = 5$  random negative pairs.

## 6.2 Parametric UMAP for Real-Time Monitoring

Train a neural network to project new experiences:

$$\mathbf{y}_{\text{new}} = f_{\theta}(\mathbf{h}_{\text{new}}) \quad (33)$$

# 7 Integration with Vector Database (Qdrant)

## 7.1 Schema Storage

Each schema is stored in the vector database with the following structure:

**Schema Representation:**

$$\mathcal{S}_k = (\mathbf{id}_k, \mathbf{z}_k, \mathcal{M}_k) \quad (34)$$

where:

- $\mathbf{id}_k \in \mathbb{N}$ : Unique schema identifier
- $\mathbf{z}_k \in \mathbb{R}^d$ : Centroid embedding vector
- $\mathcal{M}_k$ : Metadata payload containing:
  - $n_k = |C_k|$ : Number of experiences in cluster
  - $\bar{r}_k = \frac{1}{n_k} \sum_{i \in C_k} r_i$ : Average reward
  - $\sigma_k = \sqrt{\frac{1}{n_k} \sum_{i \in C_k} (r_i - \bar{r}_k)^2}$ : Volatility (std. dev. of returns)

**Storage Invariant:** All schema vectors are L2-normalized for cosine similarity search:

$$\mathbf{z}_k \leftarrow \frac{\mathbf{z}_k}{\|\mathbf{z}_k\|_2} \quad (35)$$

## 7.2 Similarity Search

Retrieve nearest schemas:

$$\mathcal{N}_k = \arg \max_k \text{cosine}(\mathbf{h}_t, \mathbf{z}_k) \quad (36)$$

## Part IV

# Neuromorphic Architecture

## Abstract

This document maps the computational components of Project JANUS to specific brain regions, ensuring biological plausibility and leveraging neuroscience insights for system design. The neuromorphic approach provides:

- **Modular Design** with clear functional boundaries
- **Biological Validation** of architectural decisions
- **Emergent Intelligence** through brain-inspired interactions

## 8 Neuromorphic Design Philosophy

### 8.1 Why Brain-Inspired Architecture?

The brain efficiently solves problems similar to trading:

- Pattern recognition under uncertainty
- Fast decision-making with delayed rewards
- Continual learning without catastrophic forgetting
- Multi-timescale memory consolidation

### 8.2 Neuroscience-to-Trading Mapping

Brain Region	Biological Function	Trading Function
Visual Cortex	Pattern recognition	GAF/ViViT chart analysis
Hippocampus	Episodic memory	Experience replay buffer
Prefrontal Cortex	Logic and planning	LTN constraint checking
Basal Ganglia	Action selection	Buy/sell/hold decisions
Cerebellum	Motor prediction	Market impact forecasting
Amygdala	Threat detection	Risk circuit breakers

## 9 Brain Region Architectures

### 9.1 Cortex: Strategic Planning & Long-term Memory

#### 9.1.1 Trading Implementation

**Component:** Neocortical Schema Network

**Implementation:**

- Schema prototypes stored in Qdrant vector database
- Each schema represents a market regime (trending, mean-reverting, volatile)
- Slow consolidation during sleep cycles

### 9.2 Hippocampus: Episodic Memory & Experience Replay

#### 9.2.1 Trading Implementation

**Component:** Episodic Buffer + SWR Replay

**Implementation:**

- Fixed-size circular buffer storing recent trades
- Sparse encoding via random projections
- Prioritized replay during training

### 9.3 Basal Ganglia: Action Selection & Reinforcement Learning

#### 9.3.1 Trading Implementation

**Component:** Dual-Pathway Decision Module

The basal ganglia implements competing pathways for action selection:

**Direct Pathway (Go Signal):**

$$\mathbf{d}_{\text{direct}} = \text{ReLU}(\mathbf{W}_{\text{direct}} \mathbf{h} + \mathbf{b}_{\text{direct}}) \quad (37)$$

where  $\mathbf{h}$  is the fused state representation and  $\mathbf{W}_{\text{direct}} \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ .

**Indirect Pathway (No-Go Signal):**

$$\mathbf{d}_{\text{indirect}} = \text{ReLU}(\mathbf{W}_{\text{indirect}} \mathbf{h} + \mathbf{b}_{\text{indirect}}) \quad (38)$$

**Action Selection:**

$$\mathbf{a}_t = \text{softmax}(\mathbf{d}_{\text{direct}} - \lambda \cdot \mathbf{d}_{\text{indirect}}) \quad (39)$$

where  $\lambda > 0$  is the inhibition weight parameter.

## 9.4 Prefrontal Cortex: Logic, Planning & Compliance

### 9.4.1 Trading Implementation

**Component:** Logic Tensor Network

Ensures regulatory compliance:

- Wash sale rules
- Position limits
- Capital allocation constraints

## 9.5 Amygdala: Fear, Threat Detection & Circuit Breakers

### 9.5.1 Trading Implementation

**Component:** Anomaly Detection Module

Triggers emergency stops based on statistical distance from normal operation:

**Mahalanobis Distance:**

$$D_M(\mathbf{s}_t) = \sqrt{(\mathbf{s}_t - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{s}_t - \boldsymbol{\mu})} \quad (40)$$

where  $\boldsymbol{\mu}$  is the historical mean state and  $\boldsymbol{\Sigma}$  is the covariance matrix.

**Circuit Breaker Condition:**

$$\text{Trigger} = \begin{cases} 1 & \text{if } D_M(\mathbf{s}_t) > \tau_{\text{danger}} \\ 0 & \text{otherwise} \end{cases} \quad (41)$$

where  $\tau_{\text{danger}}$  is calibrated to a false-positive rate (e.g.,  $\tau = 5$  for  $p < 0.001$ ).

**Additional Threat Signals:**

- Sudden volatility spike:  $\sigma_t > 3 \cdot \sigma_{\text{baseline}}$
- Drawdown threshold: cumulative loss  $> L_{\text{max}}$
- Liquidity crisis: bid-ask spread  $> 10 \times$  normal

## 9.6 Cerebellum: Motor Control & Execution

### 9.6.1 Trading Implementation

**Component:** Forward Model for Market Impact

Predicts price movement from order execution:

$$\Delta p = f_{\text{cerebellum}}(\text{order\_size}, \text{liquidity}, \text{volatility}) \quad (42)$$

## Part V

# Rust Implementation

## Abstract

This document provides production-ready Rust implementation specifications for Project JANUS, including:

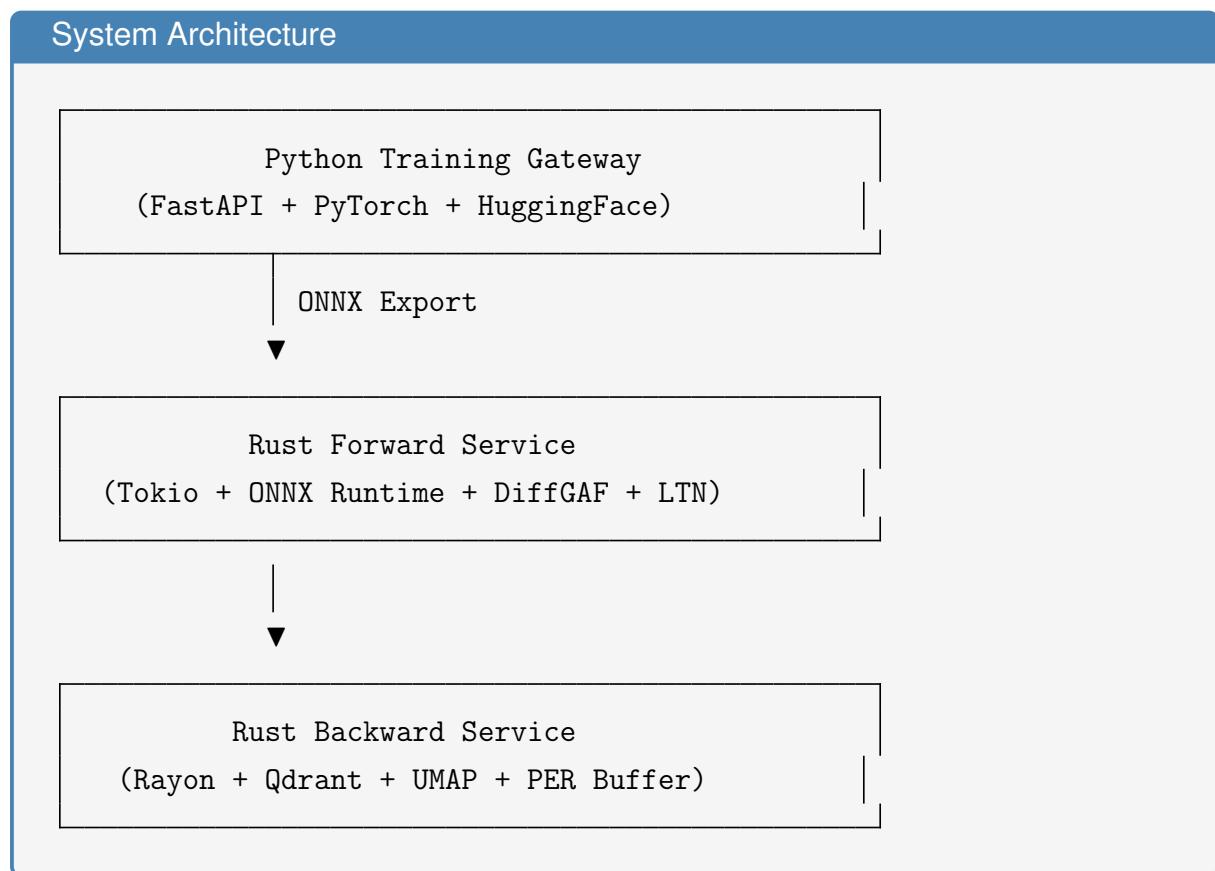
- **ML Framework Strategy** (PyTorch → ONNX → Rust inference)
- **High-Performance Services** with async Tokio runtime
- **Hybrid Training Pipeline** (Python training, Rust deployment)
- **Deployment Architecture** (Docker Compose + Kubernetes)

## 10 Architectural Overview

### 10.1 The Rust-First Philosophy

1. **Performance:** Zero-cost abstractions, no GC pauses
2. **Safety:** Memory safety without runtime overhead
3. **Concurrency:** Fearless async/await with Tokio
4. **Ecosystem:** Production-ready ML inference via ONNX

## 10.2 Component Diagram



## 11 Machine Learning Framework Strategy

### 11.1 Framework Comparison Matrix

Framework	Pros	Cons	Use Case
tch-rs	Native PyTorch, GPU support	C++ deps, larger binary	Training & inference
ONNX Runtime	Universal, production-ready	No training	Inference only
Candle	Pure Rust, HF integration	Young ecosystem	Future migration

### 11.2 Recommended Migration Path

#### 11.2.1 Phase 1: Hybrid (Months 1-3)

- Train models in PyTorch (Python)
- Export to ONNX format

- Rust inference via `ort` crate

### 11.2.2 Phase 2: Rust-Native Inference (Months 4-6)

- Optimize ONNX models for Rust
- Custom kernels for DiffGAF/LTN
- Benchmark against Python baseline

### 11.2.3 Phase 3: Full Rust ML (Months 7-12)

- Migrate training to Candle
- End-to-end Rust pipeline
- Custom GPU kernels via `wgpu`

## 12 Forward Service: Rust Implementation

### 12.1 Performance Requirements

- Latency: p99 < 10ms
- Throughput: 10,000 req/s
- Memory: < 2GB RSS

### 12.2 Core Data Structures

The system maintains several key data structures for real-time processing:

#### Market State Representation:

$$\mathcal{S}_t = (\tau_t, \mathbf{f}_t, \mathcal{O}_t, \mathbf{c}_t) \quad (43)$$

where:

- $\tau_t \in \mathbb{Z}^+$  is the timestamp
- $\mathbf{f}_t \in \mathbb{R}^d$  is the feature vector
- $\mathcal{O}_t = (\mathcal{B}_t, \mathcal{A}_t)$  is the order book with bids  $\mathcal{B}_t$  and asks  $\mathcal{A}_t$
- $\mathbf{c}_t$  contains contextual metadata (volatility, spreads, volume)

**Order Book Structure:**

$$\mathcal{B}_t = \{(p_i, q_i) : p_i \in \mathbb{R}^+, q_i \in \mathbb{R}^+\}_{i=1}^{N_{\text{bid}}} \quad (44)$$

$$\mathcal{A}_t = \{(p_j, q_j) : p_j \in \mathbb{R}^+, q_j \in \mathbb{R}^+\}_{j=1}^{N_{\text{ask}}} \quad (45)$$

## 12.3 GAF Transformation Algorithm

The GAF transformation converts time series to 2D images via the following algorithm:

**Algorithm: GAF Computation**

```

1: Input: Time series  $X = \{x_1, \dots, x_W\}$ , window size  $W$ 
2: Output: Gramian matrix  $\mathbf{G} \in \mathbb{R}^{W \times W}$ 
3:
4:  $\tilde{X} \leftarrow \text{Normalize}(X)$  to  $[-1, 1]$ 
5:  $\phi_i \leftarrow \arccos(\tilde{x}_i)$  for  $i = 1, \dots, W$ 
6: for  $i = 1$  to  $W$  do
7:   for  $j = 1$  to  $W$  do
8:      $\mathbf{G}_{ij} \leftarrow \cos(\phi_i + \phi_j)$ 
9:   end for
10: end for
11: return  $\mathbf{G}$  reshaped to  $[1, W, W]$  tensor

```

**Computational Complexity:**  $\mathcal{O}(W^2)$  for matrix construction, where  $W$  is the window size.

## 12.4 LTN Constraint Evaluation

Each constraint is represented as a weighted predicate function:

**Constraint Structure:**

$$\mathcal{C}_k = (P_k, w_k) \quad (46)$$

where  $P_k : \mathcal{S} \rightarrow [0, 1]$  is a predicate and  $w_k \in \mathbb{R}^+$  is the weight.

**Evaluation Function:**

$$\text{Eval}(\mathcal{C}_k, \mathcal{S}_t) = w_k \cdot P_k(\mathcal{S}_t) \quad (47)$$

**T-norm Operations (already defined in Part 2):**

$$a \wedge_{\mathcal{L}} b = \max(0, a + b - 1) \quad (\text{Conjunction}) \quad (48)$$

$$a \Rightarrow_{\mathcal{L}} b = \min(1, 1 - a + b) \quad (\text{Implication}) \quad (49)$$

### Total Constraint Satisfaction:

$$\mathcal{L}_{\text{constraint}} = 1 - \frac{1}{K} \sum_{k=1}^K \text{Eval}(\mathcal{C}_k, \mathcal{S}_t) \quad (50)$$

## 12.5 Async Service Architecture

The service follows an event-driven architecture with the following characteristics:

### Request Processing Pipeline:

1. **Initialization:** Load ONNX model  $\mathcal{M}_{\text{ViViT}}$  and LTN engine  $\mathcal{E}_{\text{LTN}}$
2. **Connection Handling:** Bind TCP listener on port 8080
3. **Concurrent Processing:** For each incoming request:
  - Spawn asynchronous task with model clone
  - Process request independently (non-blocking)
  - Return prediction and constraint satisfaction scores

### Concurrency Model:

$$\text{Throughput} = \frac{N_{\text{workers}} \times 1000}{T_{\text{avg}}} \quad (51)$$

where  $N_{\text{workers}}$  is the thread pool size and  $T_{\text{avg}}$  is average processing time in ms.

**Performance Characteristics:** - Non-blocking I/O via `async/await` - Zero-copy model sharing across tasks - Bounded memory through connection limiting

## 13 Backward Service: Batch Processing

### 13.1 Prioritized Experience Replay

The replay buffer maintains experiences with importance-based sampling.

#### Buffer State:

$$\mathcal{B} = \{(e_i, p_i)\}_{i=1}^N \quad (52)$$

where  $e_i$  is an experience and  $p_i \in \mathbb{R}^+$  is its priority.

#### Hyperparameters:

- $\alpha \in [0, 1]$ : Priority exponent (0 = uniform, 1 = full prioritization)
- $\beta \in [0, 1]$ : Importance sampling correction
- $C$ : Buffer capacity

**Sampling Algorithm:**

```

1: Input: Buffer  $\mathcal{B}$ , batch size  $B$ 
2: Output: Sampled batch  $\{e_{i_1}, \dots, e_{i_B}\}$ 
3:
4: Compute probabilities:  $P(i) = \frac{p_i^\alpha}{\sum_j p_j^\alpha}$ 
5: for  $k = 1$  to  $B$  do
6:   Sample index  $i_k \sim \text{Categorical}(P)$ 
7:   Add  $e_{i_k}$  to batch
8: end for
9: return batch

```

**Importance Weights:**

$$w_i = \left( \frac{1}{N \cdot P(i)} \right)^\beta \quad (53)$$

These weights correct for the non-uniform sampling distribution.

## 13.2 Schema Consolidation Algorithm

Schemas are formed by clustering experience embeddings and storing centroids.

**Algorithm: Schema Update**

```

1: Input: Experiences  $\mathcal{E} = \{e_1, \dots, e_N\}$ , number of clusters  $K$ 
2: Output: Updated schema database
3:
4: Extract embeddings:  $\mathbf{h}_i = \text{Embed}(e_i)$  for  $i = 1, \dots, N$ 
5: Cluster:  $\mathcal{C} = \{C_1, \dots, C_K\} \leftarrow \text{K-means}(\{\mathbf{h}_i\}, K)$ 
6: for  $k = 1$  to  $K$  do
7:   Compute centroid:  $\mathbf{z}_k = \frac{1}{|C_k|} \sum_{i \in C_k} \mathbf{h}_i$ 
8:   Compute statistics:
9:      $n_k = |C_k|$ 
10:     $\bar{r}_k = \frac{1}{|C_k|} \sum_{i \in C_k} r_i$  (average reward)
11:    Upsert schema  $k$  with vector  $\mathbf{z}_k$  and metadata  $(n_k, \bar{r}_k)$ 
12: end for

```

**Schema Metadata:** Each schema  $k$  stores:

- Centroid vector  $\mathbf{z}_k \in \mathbb{R}^d$
- Member count  $n_k$
- Average reward  $\bar{r}_k$
- Volatility  $\sigma_k$  (standard deviation of returns)

**K-means Objective:**

$$\min_{\mathcal{C}} \sum_{k=1}^K \sum_{i \in C_k} \|\mathbf{h}_i - \mathbf{z}_k\|^2 \quad (54)$$

## 14 Deployment Architecture

### 14.1 Service Orchestration

The system deploys as three independent services:

**Service Topology:****1. Forward Service:**

- Port: 8080 (HTTP API)
- Dependencies: ONNX model files
- Environment: Logging level, model paths
- Resource limits: 2GB memory, 2 CPU cores

**2. Backward Service:**

- Internal service (no external ports)
- Dependencies: Qdrant vector database
- Environment: Qdrant connection URL
- Scheduling: Triggered during market close

**3. Qdrant Vector Database:**

- Ports: 6333 (HTTP), 6334 (gRPC)
- Persistent storage for schemas
- Vector similarity search engine

**Service Communication:**

**Volume Management:** - Model artifacts: Shared read-only volume - Schema database: Persistent volume with backups - Experience buffer: Ephemeral storage (daily rotation)

## Conclusion

Project JANUS represents a paradigm shift in quantitative trading: from opaque black boxes to transparent, brain-inspired systems that combine the best of deep learning and symbolic reasoning.

## Key Innovations

1. **Neuromorphic Architecture:** Biologically plausible design with modular brain regions
2. **Neuro-Symbolic Fusion:** LTNs bridge neural networks and logical constraints
3. **Multi-Timescale Memory:** Three-tier hierarchy mirrors hippocampal-neocortical consolidation
4. **Production-Ready Rust:** High-performance, safe, and maintainable implementation

## Future Work

- Quantum computing integration for portfolio optimization
- Continual learning without catastrophic forgetting
- Multi-agent systems for distributed trading
- Regulatory AI for automated compliance

### Repository & Contact

**GitHub:** [https://github.com/nuniesmith/technical\\_papers](https://github.com/nuniesmith/technical_papers)

For implementation code, updates, and discussions, visit the repository.

*“The god of beginnings and transitions, looking simultaneously to the future and the past.”*