

Project JANUS

A Neuromorphic Architecture for Autonomous Trading

Technical Specification & Theoretical Foundations

*Brain-Inspired Multi-Region Architecture Integrating
Neuro-Symbolic Reasoning, Complementary Learning Systems,
and Engineered Heterogeneity*

Unified Documentation

This document consolidates all theoretical foundations and technical specifications of Project JANUS:

1. **Introduction** — Motivation, crowding resistance, and contributions
2. **Related Work** — Neuroscience-inspired AI, neuro-symbolic reasoning, and crowding literature
3. **System Architecture** — Dual-process design, brain-region mapping, and engineered heterogeneity
4. **Perception and Feature Extraction** — DiffGAF, ViViT, regime detection, and multimodal fusion
5. **Reasoning and Decision-Making** — LTN, OpAL, homeostatic regulation, and threat detection
6. **Memory and Learning** — Three-timescale memory, replay, and schema formation
7. **Implementation** — Rust-native ML stack, deployment, and safety infrastructure
8. **Limitations and Open Problems** — Honest engagement with challenges
9. **Validation Framework** — Experimental design and benchmarks

Author

Jordan Smith

Date

March 29, 2026

"The god of beginnings and transitions, looking simultaneously to the future and the past."

Abstract

Project JANUS is a neuromorphic trading architecture that maps ten functionally distinct brain regions—basal ganglia, hippocampus, cerebellum, thalamus, hypothalamus, prefrontal cortex, amygdala, visual cortex, neocortex, and central nervous system—to specialized trading subsystems within a unified autonomous agent. Grounded in the Doya–Hassabis framework of neuroscience-inspired AI (Doya, 1999; Hassabis et al., 2017), JANUS implements a dual-process architecture separating real-time perception and action (System 1, Forward Service) from offline memory consolidation and schema formation (System 2, Backward Service), following Complementary Learning Systems theory (James L McClelland, Bruce L McNaughton, and O'Reilly, 1995; Kumaran, Hassabis, and James L. McClelland, 2016).

A central design objective is *crowding resistance*: the mitigation of co-impact costs (Bucci, Mastromatteo, et al., 2020) and strategy crowding risk (Khandani and Lo, 2011; Stein, 2009) through engineered heterogeneity across deployments. Each JANUS instance is parameterized to produce idiosyncratic order flow, reducing cross-instance correlation and increasing effective market capacity (DeMiguel, Martin-Utrera, and Uppal, 2021; Wagner, 2011). The architecture brings together neuro-symbolic reasoning via Logic Tensor Networks (Badreddine et al., 2022) for differentiable constraint satisfaction, video vision transformers for visual pattern recognition (Arnab et al., 2021), Opponent Actor Learning for basal ganglia decision-making (Collins and Michael J Frank, 2014; Jaskir and Michael J. Frank, 2023), and a three-timescale memory hierarchy inspired by hippocampal sharp-wave ripple replay (M. A. Wilson and Bruce L. McNaughton, 1994; Masset et al., 2025). These subsystems differ in deployment maturity: the production trading path is a deterministic regime/rule engine, while the neuromorphic and deep-learning components that constitute the cognitive core are implemented as standalone modules not yet wired into the live inference loop (Section 10.3).

No prior published system maps multiple brain regions to distinct trading subsystems. This paper presents the theoretical foundations, mathematical specifications, neuromorphic architecture, and Rust implementation of JANUS, along with a candid assessment of limitations and a proposed validation framework.

Keywords: neuromorphic computing, algorithmic trading, neuro-symbolic AI, strategy crowding, co-impact, complementary learning systems, reinforcement learning, Rust

Contents

Abstract	2
1 Introduction	6
1.1 Motivation: Strategy Crowding and Co-Impact	6
1.2 The Quant 4.0 Paradigm	7
1.3 Contributions	7
2 Related Work	8
2.1 Neuroscience-Inspired AI	8
2.2 Complementary Learning Systems	9
2.3 Neuro-Symbolic AI	10
2.4 Neuromorphic Computing in Finance	10
2.5 Strategy Crowding Literature	10
3 System Architecture	11
3.1 Design Philosophy	11
3.2 Dual-Process Architecture	11
3.3 Brain-Region Mapping	12
3.4 Service Topology	12
3.5 Engineered Heterogeneity for Crowding Resistance	15
4 Perception and Feature Extraction	15
4.1 Differentiable Gramian Angular Fields	16
4.1.1 Normalization and Polar Encoding	16
4.1.2 Gramian Field Generation	17
4.1.3 Differentiable Backpropagation	17
4.1.4 Spatiotemporal GAF Video	17
4.2 Vision Model Architecture	17
4.2.1 DiffGAF-LSTM	18
4.2.2 Video Vision Transformer (ViViT)	18
4.3 Ensemble Regime Detection	19
4.3.1 Detection Methods	19
4.3.2 Regime Categories	19
4.4 Multimodal Fusion via Thalamic Gating	20
4.4.1 Input Modalities	20
4.4.2 Gated Cross-Attention	20
5 Reasoning and Decision-Making	21
5.1 Neuro-Symbolic Constraints via Logic Tensor Networks	21

5.1.1	Mathematical Foundation	21
5.1.2	Fuzzy Logic Operations	22
5.1.3	Knowledge Base and Constraints	23
5.1.4	Logical Loss Function	23
5.2	Action Selection via Opponent Actor Learning	24
5.2.1	Dual-Pathway Architecture	24
5.2.2	Brain Wiring Pipeline	25
5.3	Homeostatic Risk Regulation	26
5.4	Threat Detection and Circuit Breakers	26
5.4.1	Anomaly Detection	26
5.4.2	Flow Toxicity and Circuit Breakers	27
5.4.3	Fear Extinction	28
5.5	Execution via Cerebellar Forward Models	28
5.5.1	Market Impact Prediction	28
5.5.2	Execution Error Correction	28
6	Memory and Learning	28
6.1	Three-Timescale Memory Hierarchy	29
6.2	Short-Term: Hippocampal Episodic Buffer	29
6.2.1	Episodic Buffer	29
6.2.2	Pattern Separation and Spatial Mapping	30
6.3	Medium-Term: Sharp-Wave Ripple Replay	30
6.3.1	Replay Prioritization	30
6.3.2	Sleep-Phase Consolidation	31
6.4	Long-Term: Neocortical Schema Formation	31
6.4.1	Schema Representation	31
6.4.2	Recall-Gated Consolidation	32
6.5	UMAP Visualization and Drift Detection	33
7	Implementation	33
7.1	Rust-Native ML Stack	34
7.2	Service Architecture and Deployment	35
7.2.1	Service Topology	35
7.2.2	Execution Service	37
7.3	Trading Strategies and Regime-Aware Gating	37
7.4	Safety Infrastructure	38
7.4.1	Preflight Validation	39
7.4.2	Runtime Monitoring	39
8	Limitations and Open Problems	40

8.1	Research-to-Production Maturity Gap	40
8.2	The Brain Metaphor Critique	40
8.3	Integration Complexity	41
8.4	Non-Stationarity and Regime Shifts	41
8.5	Backtest Overfitting Risk	42
8.6	Crowding Resistance Is Not Crowding Immunity	42
8.7	LTN Scalability	42
8.8	Adversarial Robustness	43
9	Validation Framework	43
9.1	Multi-Agent LOB Simulation	43
9.2	Walk-Forward Backtesting	44
9.3	Ablation Studies	44
9.4	Benchmark Comparisons	44
10	Conclusion	45
10.1	Key Innovations	45
10.2	Future Work	46
10.3	Implementation Status	46

1 Introduction

1.1 Motivation: Strategy Crowding and Co-Impact

Strategy crowding—the convergence of multiple algorithmic trading systems on similar positions—poses systemic risk through amplified co-impact costs and correlated drawdowns.

The *square-root impact law*—stating that the price impact of a metaorder scales as $\Delta p \sim \sigma \sqrt{Q/V}$, where Q is order volume and V is daily volume—is one of the most robust empirical regularities in market microstructure (Tóth et al., 2011; Bouchaud et al., 2018; Bucci, Benzaquen, et al., 2019). Bucci, Mastromatteo, et al. (2020) extended this framework to show how *simultaneous* institutional metaorders interact through net order flow: when N agents each trade Q/N , the market responds to aggregate flow without distinguishing individual metaorders. Critically, co-impact introduces a finite intercept I_0 that grows with sign correlation among agents—meaning correlated (crowded) trading amplifies impact costs beyond what single-agent models predict.

The canonical empirical demonstration is the August 2007 quant meltdown (Khandani and Lo, 2011), where quantitative equity funds that had developed models independently suffered coordinated losses because their strategies had converged on similar factor exposures. Stein (2009) formalized this as a coordination problem: individual traders cannot observe how many others pursue the same strategy, creating negative externalities through crowded-trade effects and leverage decisions.

Subsequent empirical work has confirmed crowding’s destructive potential. Kralingen et al. (2021) showed that market clustering—measured against a maximum-entropy null model—has a *causal* effect on tail risk, even after controlling for standard risk drivers. Brown, Howard, and Lundblad (2022) demonstrated that crowded stocks outperform non-crowded stocks on average but with substantially elevated tail risk. Lou and Polk (2022) proposed “comomentum” (high-frequency abnormal return correlation among momentum stocks) as a crowding measure, finding sharp reversals when comomentum is high.

Contemporary algorithmic trading systems are particularly susceptible to crowding. As firms adopt similar pre-trained foundation models, output homogeneity increases regardless of nominal strategy differences. Baltas (2019) classifies alternative risk premia (ARP) into *divergence premia* (e.g., momentum—self-reinforcing, no fundamental anchor) and *convergence premia* (e.g., value—self-correcting), showing that divergence premia underperform following crowded periods; Volpati et al. (2020) identifies significant and increasing crowding in Fama–French factors. Furthermore, Stillman and Baggott (2024) demonstrates that groups of homogeneous neuro-symbolic traders in virtual markets produce *price suppression*, directly illustrating the risks of AI crowd

homogeneity. JANUS’s multi-region, heterogeneous architecture is designed to resist these effects, as detailed in Section 3.5.

1.2 The Quant 4.0 Paradigm

The trajectory of algorithmic trading has historically been defined by a tension between interpretability and capability. We are currently witnessing a phase transition from the “black box” empiricism of deep learning (LeCun, Bengio, and Hinton, 2015) toward a new paradigm of Neuro-Symbolic integration (Marra et al., 2024; A. d. Garcez and Lamb, 2023). Project JANUS stands at the vanguard of this transition, termed **Quant 4.0**. This architecture reimagines the financial agent as a biological entity—one that perceives, reasons, remembers, and fears.

The intellectual foundation rests on three convergent programs: (i) neuroscience-inspired AI (Hassabis et al., 2017), which argues that understanding the computational principles of the brain remains the richest source of algorithmic innovation; (ii) the Adaptive Markets Hypothesis (Lo, 2004; Lo, 2017), which replaces the Efficient Markets assumption with an evolutionary framework where market agents adapt, compete, and are selected; and (iii) neuro-symbolic reasoning (A. d. Garcez and Lamb, 2023; Badreddine et al., 2022), which bridges the gap between statistical pattern recognition and logical constraint satisfaction.

Historical Evolution of Quantitative Finance

- **Quant 1.0 (1980s–1990s):** Era of heuristics and expert systems with high interpretability but extreme rigidity
- **Quant 2.0 (1990s–2000s):** Statistical rigour through mean reversion, cointegration, and factor models (Fama and French, 1993)
- **Quant 3.0 (2010s–Present):** Deep learning hegemony with LSTMs, Transformers (Vaswani et al., 2017), and deep reinforcement learning (Sutton and Barto, 2018)
- **Quant 4.0 (JANUS):** Neuro-Symbolic AI achieving adaptability of deep learning with reliability of rule-based systems, grounded in neuroscience

1.3 Contributions

The principal contributions of this work are:

1. **Neuromorphic architecture:** The first trading system to map ten functionally distinct brain regions to specialized subsystems, grounded in the Doya–Hassabis

framework (Doya, 1999; Hassabis et al., 2017) and the super-learning hypothesis (Caligiore et al., 2019).

2. **Crowding resistance via engineered heterogeneity:** A systematic approach to reducing co-impact costs through parameterized diversity across deployments (Wagner, 2011; DeMiguel, Martin-Utrera, and Uppal, 2021).
3. **Neuro-symbolic fusion:** The first application of Logic Tensor Networks (Badreddine et al., 2022) to financial decision-making, formulating regulatory and risk constraints as differentiable fuzzy-logic axioms.
4. **Opponent Actor Learning for trading:** A design for financial action selection based on basal ganglia computational models (Collins and Michael J Frank, 2014; Jaskir and Michael J. Frank, 2023)—to our knowledge the first such application to trading.
5. **Three-timescale memory:** A Complementary Learning Systems architecture (James L McClelland, Bruce L McNaughton, and O'Reilly, 1995; Kumaran, Hassabis, and James L. McClelland, 2016) with biologically inspired sharp-wave ripple replay.
6. **Production Rust implementation:** An end-to-end Rust-native ML stack with eight-service deployment architecture.

The remainder of this paper is organized as follows. Section 2 surveys the relevant literature. Section 3 presents the system architecture. Sections 4–6 detail the three functional layers: perception, reasoning, and memory. Section 7 covers the Rust implementation. Section 8 discusses limitations, and Section 9 describes the validation framework.

2 Related Work

2.1 Neuroscience-Inspired AI

JANUS's multi-region architecture applies functional brain-region decomposition to trading system design, grounded in the Doya–Hassabis framework of neuroscience-inspired AI (Doya, 1999; Hassabis et al., 2017).

Doya (1999) proposed the canonical functional decomposition of brain regions by learning type:

- **Cerebellum** → supervised learning (forward models, error correction)
- **Basal ganglia** → reinforcement learning (reward-driven action selection)

- **Cerebral cortex** → unsupervised learning (schema formation, pattern discovery)

Doya (2002) extended this framework to show that neuromodulators (dopamine, serotonin, norepinephrine, acetylcholine) regulate meta-parameters of learning—learning rate, discount factor, exploration–exploitation balance—providing a biologically grounded mechanism for adaptive parameter control. Caligiore et al. (2019) formalized the “super-learning hypothesis,” demonstrating that the integration of learning processes *across* cortex, cerebellum, and basal ganglia produces capabilities exceeding any single region.

Hassabis et al. (2017) articulated the programmatic statement for neuroscience-inspired AI: the goal is *functional* inspiration, not biological simulation. JANUS adopts this position, with brain-region mappings (Table 1) justified by functional analogy—each region implements a computational principle (e.g., supervised error correction, reinforcement-driven action selection, episodic memory replay) that addresses a specific trading requirement.

This approach is further validated by Yamakawa (2021), who proposed the “whole brain architecture” approach as a systematic method for developing artificial general intelligence by referencing brain structure, and by Caligiore et al. (2017), who established consensus on the functional interactions between cerebellum, basal ganglia, and cortex.

2.2 Complementary Learning Systems

The separation of JANUS into Forward (System 1) and Backward (System 2) services directly implements Complementary Learning Systems (CLS) theory (James L McClelland, Bruce L McNaughton, and O’Reilly, 1995). Kumaran, Hassabis, and James L. McClelland (2016) updated CLS theory (“CLS 2.0”), establishing that intelligent agents require both a fast-learning system (hippocampus, for rapid encoding of specific experiences) and a slow-learning system (neocortex, for gradual extraction of statistical regularities). JANUS maps these to:

- **Forward Service (hippocampal analogue):** Rapid encoding of each trading episode with episodic specificity
- **Backward Service (neocortical analogue):** Slow consolidation of episodic memories into generalized market schemas through sharp-wave ripple replay (M. A. Wilson and Bruce L. McNaughton, 1994; Buzsáki, 2015)

Masset et al. (2025) recently demonstrated that dopaminergic neurons encode reward prediction errors at diverse temporal discount factors, providing direct biological validation for JANUS’s three-timescale memory architecture (short-term hippocampal buffer, medium-term SWR consolidation, long-term neocortical schemas).

2.3 Neuro-Symbolic AI

Logic Tensor Networks (LTNs) (Badreddine et al., 2022; Serafini and A. S. d. Garcez, 2016) bridge neural networks and first-order logic by grounding logical symbols in real-valued tensors and replacing Boolean connectives with differentiable fuzzy operators (Łukasiewicz t-norms). This enables gradient-based optimization of logical satisfiability, allowing domain knowledge expressed as first-order axioms to constrain and guide neural learning. Marra et al. (2024) and A. d. Garcez and Lamb (2023) survey the broader neuro-symbolic landscape, establishing that hybrid architectures combining sub-symbolic pattern recognition with symbolic constraint satisfaction consistently outperform purely neural approaches in domains requiring safety guarantees, regulatory compliance, and interpretable reasoning—precisely the requirements of autonomous trading. No prior work applies LTNs to financial decision-making; this represents a novel contribution of JANUS.

2.4 Neuromorphic Computing in Finance

The intersection of neuromorphic computing and finance is extremely sparse. Mohan et al. (2025) apply spiking neural networks to cross-market portfolio optimization but use no brain-region mapping. Bi and Calhoun (2025) propose a “Financial Connectome” using brain connectivity analysis methods on market data, but this applies neuroscience *analysis tools* to markets rather than neuroscience *architectural principles* to trading systems. Ezinwoke and Rhodes (2025) apply SNNs to HFT price prediction without multi-region architecture. **No published system maps multiple brain regions to distinct trading subsystems.** JANUS is the first.

2.5 Strategy Crowding Literature

The empirical foundations of strategy crowding rest on the square-root impact law (Tóth et al., 2011; Bouchaud et al., 2018; Bucci, Benzaquen, et al., 2019) and its multi-agent extension through co-impact (Bucci, Mastromatteo, et al., 2020). Stein (2009) formalized the theoretical coordination failure: individual agents cannot observe aggregate strategy overlap, creating negative externalities. Khandani and Lo (2011) provided the canonical case study with the August 2007 quant meltdown, demonstrating that independently developed strategies can converge on similar factor exposures. More recent empirical work by Kralingen et al. (2021) establishes a causal link between market clustering and tail risk, while Caccioli et al. (2014) analyzes the systemic stability implications of crowded portfolios.

On the mitigation side, Wagner (2011) shows that rational agents should choose heterogeneous portfolios to avoid joint liquidation risk, and DeMiguel, Martin-Utrera,

and Uppal (2021) demonstrates that “trading diversification”—institutions exploiting different characteristics—increases effective market capacity by 45%. Baltas (2019) distinguishes divergence and convergence premia in the context of alternative risk premia crowding, and Volpati et al. (2020) identifies increasing crowding in standard Fama–French factors. These results collectively motivate JANUS’s engineered heterogeneity approach (Section 3.5).

3 System Architecture

3.1 Design Philosophy

JANUS adopts a brain-inspired modular architecture in which each subsystem implements the computational principle of a specific brain region, following the Doya–Hassabis framework (Doya, 1999; Hassabis et al., 2017). The brain efficiently solves problems structurally similar to trading (Nathaniel D Daw et al., 2006), demonstrating capabilities that map directly to trading challenges:

- Pattern recognition under uncertainty (visual cortex and hippocampus (Buzsáki, 2015))
- Fast decision-making with delayed rewards (basal ganglia (Collins and Michael J Frank, 2014))
- Continual learning without catastrophic forgetting (complementary learning systems (James L McClelland, Bruce L McNaughton, and O’Reilly, 1995))
- Multi-timescale memory consolidation (hippocampus to neocortex (Buzsáki, 2015))
- Homeostatic regulation under varying conditions (hypothalamic control (Sterling, 2012))

Rather than biological simulation, JANUS pursues *functional* inspiration: each brain region is implemented as a software module whose computational role mirrors that of its biological counterpart. This modularity enables independent development, testing, and deployment of each region while preserving the emergent intelligence that arises from their interaction (Caligiore et al., 2019).

3.2 Dual-Process Architecture

JANUS separates real-time perception and action from offline memory consolidation, implementing Dual-Process Theory (Kahneman, 2011; J. S. B. Evans, 2008; J. S. B. T.

Evans and Stanovich, 2013) as a concrete engineering design. The system divides into two complementary services:

Service	Persona	Cognitive Role	Biological Analogue
Forward Service	Janus Bifrons	Perception & Action (System 1)	Basal Ganglia & Thalamus
Backward Service	Janus Consivius	Memory & Learning (System 2)	Hippocampus & Neocortex

This separation allows JANUS to optimize for latency on the hot path (Forward Service) while reserving heavy computational resources for consolidation and schema formation on the cold path (Backward Service), implementing CLS theory (James L McClelland, Bruce L McNaughton, and O'Reilly, 1995; Kumaran, Hassabis, and James L. McClelland, 2016).

In addition to the core Forward and Backward services, the production system deploys three supporting services: an **Execution Service** for multi-exchange order routing, a **Data Service** for centralized market data management, and a **CNS (Central Nervous System) Service** for system-wide health monitoring and preflight validation (Section 7.4).

3.3 Brain-Region Mapping

This mapping is grounded in empirical neuroscience and cognitive modeling (Michael J Frank, Loughry, and O'Reilly, 2006; Collins and Michael J Frank, 2014; Foster, Morris, and Dayan, 2013). JANUS implements ten brain regions, each serving a distinct functional role:

Each brain region is detailed in subsequent sections: visual cortex and thalamus in Section 4, basal ganglia and prefrontal cortex in Section 5, and hippocampus and cortex in Section 6.


3.4 Service Topology

The production deployment comprises eight interconnected services organized in a layered topology:

Table 1: Neuroscience-to-trading mapping for JANUS brain regions

Brain Region	Biological Function	Trading Function
Visual Cortex	Pattern recognition	GAF/ViViT chart analysis (Wang and Oates, 2015; Arnab et al., 2021)
Hippocampus	Episodic memory (Buzsáki, 2015)	Experience replay buffer (Schaul et al., 2015)
Prefrontal Cortex	Logic and planning (Michael J Frank, Loughry, and O'Reilly, 2006)	LTN constraint checking (Badreddine et al., 2022)
Basal Ganglia	Action selection (Collins and Michael J Frank, 2014)	Buy/sell/hold decisions
Cerebellum	Motor prediction (Wolpert, Miall, and Kawato, 1998)	Market impact forecasting (Almgren and Chriss, 2001)
Amygdala	Threat detection (LeDoux, 2000)	Risk circuit breakers
Thalamus	Attentional gating (Halassa and Kastner, 2017)	Modality fusion and signal routing
Hypothalamus	Homeostatic regulation (Sterling, 2012)	Position sizing and risk appetite
Cortex	Strategic planning	Regime schemas and hierarchical RL
Integration	Inter-region coordination	Service bridges and data pipeline

System Architecture (Pure Rust)



The Forward Service acts as the primary orchestrator during market hours, routing signals through the neuromorphic brain pipeline (regime → hypothalamus → amygdala → gating → correlation → execution) and dispatching orders to the Execution Service. The Backward Service operates asynchronously, consuming episodic memories from Redis and consolidating them into long-term schemas stored in PostgreSQL. The CNS Service continuously monitors all components, enforcing health invariants and executing preflight validation before each trading session.

3.5 Engineered Heterogeneity for Crowding Resistance

JANUS addresses crowding through *engineered heterogeneity*: each deployment is configured to produce idiosyncratic order flow that minimizes cross-instance correlation. This is grounded in two key theoretical results:

- Wagner (2011): Rational agents should choose heterogeneous portfolios and forgo diversification benefits to avoid joint liquidation risk. Portfolio heterogeneity directly reduces systemic risk.
- DeMiguel, Martin-Utrera, and Uppal (2021): “Trading diversification”—institutions exploiting different characteristics—increases capacity by 45%, optimal investment by 43%, and profits by 22%.

The heterogeneity is implemented across multiple brain regions:

Table 2: Sources of engineered heterogeneity across brain regions

Brain Region	Heterogeneity Mechanism	Effect on Order Flow
Thalamus	Modality attention weights randomized at initialization	Different signal weighting per instance
Hypothalamus	Risk setpoints drawn from individualized distributions	Different position sizing and risk tolerance
Basal Ganglia	Dopamine sensitivity and discount factors varied	Different action selection thresholds
Prefrontal	LTN axiom weights and hedge parameters varied	Different constraint priorities
Hippocampus	Replay priorities and consolidation schedules varied	Different learning from the same experience
Cerebellum	Execution timing models calibrated to different horizons	Different order placement patterns

Important caveat: The literature supports heterogeneity as a mechanism for crowding *resistance*, not crowding *immunity*. The 2007 quant quake demonstrated that independently developed strategies can converge on similar factor exposures despite nominal diversity (Khandani and Lo, 2011). Evolutionary dynamics research shows that crowding effects can emerge even in heterogeneous populations when memory sizes are small. We therefore claim crowding resistance, not elimination, and note that empirical validation via multi-agent simulation is required (see Section 9).

4 Perception and Feature Extraction

JANUS’s sensory layer—implemented through the *Visual Cortex* and *Thalamus* brain regions—transforms raw market data into rich perceptual representations that drive

all downstream reasoning. This section describes the full perception pipeline: visual encoding via Gramian Angular Fields, vision model architectures, ensemble regime detection, and multimodal fusion through thalamic gating.

4.1 Differentiable Gramian Angular Fields

The visual subsystem transforms time series data into spatiotemporal images, enabling the system to “see” market patterns that traditional numerical methods miss. This approach is grounded in the work of Wang and Oates (2015), who demonstrated that imaging time series significantly improves classification and imputation tasks by exposing temporal correlations to the inductive biases of convolutional neural networks. Research by Chen and Tsai (2020) has validated that GAF-based encodings of financial candlestick data substantially outperform raw time-series inputs in classification tasks by capturing multi-scale temporal structures.

Given raw market data $X = \{x_1, x_2, \dots, x_T\}$ where $x_t \in \mathbb{R}^D$ (multi-feature time series), we first apply feature selection to extract F relevant features, then encode the selected features through the three-step pipeline below.

4.1.1 Normalization and Polar Encoding

Step 1: Normalization. For inference, we apply min-max normalization to the domain $[-1, 1]$, followed by Piecewise Aggregate Approximation (PAA) for temporal resizing:

$$\tilde{x}_t = 2 \cdot \frac{x_t - x_{\min}}{x_{\max} - x_{\min}} - 1 \quad (1)$$

ensuring the subsequent $\arccos$ operation is well-defined with $\tilde{x}_t \in [-1, 1]$.

For the training pipeline, we employ learnable affine transformations with domain constraints:

$$\tilde{x}_t = \tanh\left(\gamma \odot \frac{x_t - \mu}{\sigma} + \beta\right) \quad (2)$$

where $\gamma, \beta \in \mathbb{R}^F$ are learned parameters, and μ, σ are running statistics. The $\tanh$ function guarantees $\tilde{x}_t \in (-1, 1)$.

Step 2: Polar Coordinate Transformation. Map normalized values to angular space:

$$\phi_t = \arccos(\tilde{x}_t) \in [0, \pi] \quad (3)$$

$$r_t = \frac{t}{T} \quad (\text{normalized timestamp}) \quad (4)$$

4.1.2 Gramian Field Generation

Construct the Gramian Angular Summation Field (GASF):

$$\mathbf{G}_{ij} = \cos(\phi_i + \phi_j) = \tilde{x}_i \tilde{x}_j - \sqrt{1 - \tilde{x}_i^2} \sqrt{1 - \tilde{x}_j^2} \quad (5)$$

Or the Gramian Angular Difference Field (GADF), which JANUS employs to encode velocity of price changes as visual textures:

$$\mathbf{G}_{ij} = \sin(\phi_i - \phi_j) = \sqrt{1 - \tilde{x}_i^2} \tilde{x}_j - \tilde{x}_i \sqrt{1 - \tilde{x}_j^2} \quad (6)$$

This transformation allows the system to visually perceive volatility regimes and microstructure dynamics (Wang and Oates, 2015).

4.1.3 Differentiable Backpropagation

The Rust implementation provides a fully differentiable GAF engine with analytically computed Jacobians for end-to-end gradient flow. The key derivative through the polar mapping is:

$$\frac{d\phi_k}{d\tilde{x}_k} = \frac{-1}{\sqrt{1 - \tilde{x}_k^2}} \quad (7)$$

which enables backpropagation through the entire GAF encoding pipeline. Numerical gradient verification tests confirm correctness within 10^{-4} tolerance.

4.1.4 Spatiotemporal GAF Video

To capture temporal dynamics, we generate a sequence of GAF frames using sliding windows. Given a time series of length T , window size W , and stride S :

1. Extract windows: $X_k = \{x_{(k-1)S+1}, \dots, x_{(k-1)S+W}\}$ for $k = 1, \dots, N$
2. Generate GAF for each window: $\mathbf{G}_k = \text{GAF}(X_k) \in \mathbb{R}^{W \times W}$
3. Stack into video: $\mathbf{V} = [\mathbf{G}_1, \mathbf{G}_2, \dots, \mathbf{G}_N] \in \mathbb{R}^{N \times W \times W}$

The data ingestion pipeline handles windowed buffering and streaming through dedicated preprocessing modules.

Computational Complexity: $\mathcal{O}(W^2)$ for matrix construction, where W is the window size. For the full GAF video of N windows this becomes $\mathcal{O}(NW^2)$, computed in parallel across windows using the Tokio async runtime.

4.2 Vision Model Architecture

Algorithm 1 GAF Computation

```

1: Input: Time series  $X = \{x_1, \dots, x_W\}$ , window size  $W$ 
2: Output: Gramian matrix  $\mathbf{G} \in \mathbb{R}^{W \times W}$ 
3:
4:  $\tilde{X} \leftarrow \text{Normalize}(X)$  to  $[-1, 1]$ 
5:  $\phi_i \leftarrow \arccos(\tilde{x}_i)$  for  $i = 1, \dots, W$ 
6: for  $i = 1$  to  $W$  do
7:   for  $j = 1$  to  $W$  do
8:      $\mathbf{G}_{ij} \leftarrow \cos(\phi_i + \phi_j)$ 
9:   end for
10: end for
11: return  $\mathbf{G}$  reshaped to  $[1, W, W]$  tensor

```

4.2.1 DiffGAF-LSTM

The DiffGAF-LSTM vision model pairs the DiffGAF encoding with an LSTM temporal model. GAF frames are encoded and fed sequentially into an LSTM network that captures inter-frame temporal dependencies. This architecture provides robust performance while maintaining the differentiability required for end-to-end training.

4.2.2 Video Vision Transformer (ViViT)

The ViViT model uses a factorized spatiotemporal transformer (Arnab et al., 2021). Unlike standard Vision Transformers (Dosovitskiy et al., 2020) which process static images, ViViT factorizes attention across both space (price/volume levels of the limit order book) and time (sequences of LOB snapshots), enabling the system to track dynamic microstructure events.

Patch Embedding: Divide each frame $\mathbf{G}_k$ into non-overlapping patches:

$$\mathbf{P}_k = \text{Reshape}(\mathbf{G}_k) \in \mathbb{R}^{P \times (p^2)} \quad (8)$$

where $P = (W/p)^2$ is the number of patches per frame.

Spatial Attention: Apply self-attention within each frame:

$$\mathbf{Z}_k^{(l)} = \text{MSA}(\text{LN}(\mathbf{Z}_k^{(l-1)})) + \mathbf{Z}_k^{(l-1)} \quad (9)$$

Temporal Attention: Apply attention across frames:

$$\mathbf{H}^{(l)} = \text{MSA}(\text{LN}([\mathbf{Z}_1^{(l)}, \dots, \mathbf{Z}_N^{(l)}])) \quad (10)$$

Both the DiffGAF-LSTM and ViViT architectures are implemented natively in Rust, with the system selecting the appropriate model based on configuration and available resources. The ViViT model gracefully degrades to the DiffGAF-LSTM model when

computational constraints require it.

4.3 Ensemble Regime Detection

Market regime identification is a critical upstream component that conditions all downstream decision-making. JANUS employs an ensemble approach combining multiple detection methods to robustly classify market states.

4.3.1 Detection Methods

Hidden Markov Model. A Gaussian HMM models latent regime states with observable market features:

$$P(\mathbf{x}_t \mid z_t = k) = \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k) \quad (11)$$

where $z_t \in \{1, \dots, K\}$ is the latent regime state and the transition matrix $\mathbf{A}_{ij} = P(z_{t+1} = j \mid z_t = i)$ captures regime persistence and switching dynamics.

Statistical Methods. Complementing the HMM, JANUS applies rolling statistical tests including variance ratio tests for mean-reversion detection, Hurst exponent estimation for trend persistence, and distributional tests for fat-tail identification.

Technical Methods. Classical technical analysis signals (trend strength indicators, volatility measures, momentum oscillators) are fused with the statistical methods to provide corroborating evidence for regime classification.

4.3.2 Regime Categories

The ensemble classifier identifies seven market regimes, each mapped to distinct trading behavior profiles:

Regime	Characteristics	System Behavior
Bull	Sustained upward trend, expanding volume	Amplify Direct (Go) pathway
Bear	Sustained downward trend, risk-off	Amplify Indirect (No-Go) pathway
Ranging	Low directional conviction, bounded action	Favor mean-reversion strategies
Crisis	Extreme volatility, correlation breakdown	Engage Amygdala circuit breakers
Recovery	Post-crisis normalization, lower volatility	Gradually release risk constraints
Bubble	Parabolic acceleration, euphoric sentiment	Heighten caution via Hypothalamus
Deflation	Sustained contraction, liquidity withdrawal	Reduce position sizing

Table 3: JANUS regime categories and their associated system behaviors.

The detected regime feeds into the Hypothalamus module for homeostatic regulation and the Basal Ganglia for dopamine-modulated action selection, forming the first stage of the brain wiring pipeline.

4.4 Multimodal Fusion via Thalamic Gating

JANUS integrates multiple data modalities through gated cross-attention (Alayrac et al., 2022), allowing the system to dynamically weight inputs based on their predictive uncertainty. This is the machine-learning analogue of thalamic attentional gating in biological systems (Halassa and Kastner, 2017; George et al., 2025).

4.4.1 Input Modalities

The production system processes four specialized data streams through dedicated fusion engines:

- **Order Book:** Full limit order book depth, bid-ask dynamics, and microstructure features
- **Price:** Multi-timeframe price action including OHLCV and derived indicators, with Chronos foundation model forecasting (Ansari et al., 2024b; Ansari et al., 2024a)
- **Volume:** Volume profile analysis, volume-weighted metrics, and participation rates
- **Sentiment:** BERT/FinBERT embeddings (Hugging Face, 2024) from news feeds (NewsAPI, CryptoPanic) and social sentiment signals, with Qdrant-backed similarity search for regime-aware aggregation

Additionally, the system integrates supplementary data sources including weather data (via OpenWeatherMap) and space weather/celestial data for correlation analysis with commodity and energy markets.

4.4.2 Gated Cross-Attention

Attention Computation. Following the attention mechanism of Vaswani et al. (2017), with support for causal masking and multi-head configuration:

$$\text{Attn}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^T}{\sqrt{d_k}}\right) \mathbf{V} \quad (12)$$

Gating Mechanism. The gating function suppresses noise and amplifies signal, similar to thalamic regulation (Halassa and Kastner, 2017):

$$g = \text{sigmoid}(\mathbf{W}_g[\mathbf{v}; \mathbf{t}; \mathbf{s}] + \mathbf{b}_g) \quad (13)$$

Each modality-specific fusion engine applies this gating independently, allowing the system to dynamically suppress noisy order book data during low-liquidity periods while amplifying sentiment signals during news-driven regimes.

The Thalamic Reticular Nucleus further enriches attentional control through oscillatory dynamics: Wilson-Cowan mean-field models (H. R. Wilson and Cowan, 1972; Y. Zhang et al., 2024) implement coupled excitatory-inhibitory neural populations via ODEs with Hilbert-transform-based amplitude/phase estimation, bifurcation analysis, neuromodulator-driven regime switching, and fixed-point stability classification. These oscillatory dynamics enable JANUS to modulate attentional focus on timescales aligned with market microstructure rhythms.

5 Reasoning and Decision-Making

This section describes JANUS’s cognitive core: the brain regions responsible for logical reasoning, action selection, risk regulation, threat detection, and execution planning. Together, these modules form the decision-making pipeline that transforms perceptual representations (Section 4) into concrete trading actions while enforcing regulatory compliance, managing risk homeostatically, and maintaining robust circuit-breaker safeguards.

5.1 Neuro-Symbolic Constraints via Logic Tensor Networks

LTNs bridge neural networks and first-order logic, enabling differentiable constraint satisfaction (Badreddine et al., 2022; Serafini and A. S. d. Garcez, 2016). This neuro-symbolic approach allows JANUS to enforce regulatory and risk constraints directly within gradient descent optimization, a capability absent in pure deep learning systems (Marra et al., 2024; A. d. Garcez and Lamb, 2023). No prior work applies LTN to financial decision-making; this represents a novel contribution of JANUS.

5.1.1 Mathematical Foundation

The central innovation of LTNs (Badreddine et al., 2022) is the ability to make Boolean logic differentiable using Real Logic, specifically Łukasiewicz t-norms.

Grounding Function. Map logical constants to real vectors:

$$\mathcal{G} : \mathcal{C} \rightarrow \mathbb{R}^d \tag{14}$$

Predicate Grounding. A predicate $P(x)$ is grounded as a neural network $f_\theta : \mathbb{R}^d \rightarrow [0, 1]$, where truth values range continuously from 0 to 1 rather than being discrete

binary values.

5.1.2 Fuzzy Logic Operations

Following Badreddine et al. (2022), we employ fuzzy logic operators based on Łukasiewicz t-norms that are differentiable and thus compatible with backpropagation.

Conjunction (AND). For training, we use Product Logic to ensure smooth gradients:

$$u \wedge v = u \cdot v \quad (15)$$

For inference/evaluation, standard Łukasiewicz logic is used:

$$u \wedge v = \max(0, u + v - 1) \quad (16)$$

Disjunction (OR).

$$u \vee v = \min(1, u + v) \quad (17)$$

Negation (NOT).

$$\neg u = 1 - u \quad (18)$$

Implication (IF-THEN). For training (Product Logic):

$$u \Rightarrow v = 1 - u + u \cdot v \quad (19)$$

For inference (Łukasiewicz Logic):

$$u \Rightarrow v = \min(1, 1 - u + v) \quad (20)$$

Bi-Implication (Equivalence).

$$u \Leftrightarrow v = 1 - |u - v| \quad (21)$$

Linguistic Hedges. The implementation extends classical fuzzy logic with linguistic hedge operators for nuanced truth-value modification:

$$\text{very}(x) = x^2 \quad (22)$$

$$\text{somewhat}(x) = \sqrt{x} \quad (23)$$

$$\text{slightly}(x) = \sqrt{x} - x \quad (24)$$

$$\text{extremely}(x) = x^3 \quad (25)$$

These hedges allow more expressive constraint formulation (e.g., “very risky” vs. “somewhat risky” conditions).

5.1.3 Knowledge Base and Constraints

The knowledge base $\mathcal{KB}$ encodes regulatory constraints and risk management rules as logical predicates. The implementation provides a comprehensive compliance engine covering multiple constraint categories.

Wash Sale Constraint. The Wash Sale Rule (Internal Revenue Service, 2024)—a critical regulatory constraint for active traders—prevents claiming tax losses on securities sold and repurchased within 30 days. The implementation enforces the full 30-day window both before and after the sale, tracks loss sales per symbol, computes disallowed loss amounts, and blocks violating trades:

$$\forall t : \text{Sell}(t) \wedge \text{Buy}(t') \wedge |t - t'| < 30 \Rightarrow \neg \text{TaxLoss}(t) \quad (26)$$

Almgren-Chriss Risk Constraint. Following the optimal execution framework of Almgren and Chriss (2001), we constrain market impact relative to volatility:

$$\forall \text{order} : \text{Execute}(\text{order}) \Rightarrow \text{Slippage}(\text{order}) < \lambda \cdot \text{Volatility} \quad (27)$$

This ensures trades remain within the efficient frontier between expected cost and risk.

Additional Constraint Categories. Beyond the foundational constraints above, the production knowledge base includes:

- **Position Limits:** Maximum exposure per asset and aggregate portfolio
- **Capital Allocation:** Constraints on capital deployment across strategies
- **Risk Limits:** Value-at-Risk, maximum drawdown, and volatility thresholds
- **Proprietary Firm Rules:** Compliance with specific prop trading firm requirements (daily loss limits, trailing drawdown, consistency rules)

5.1.4 Logical Loss Function

Satisfiability Aggregation.

$$\text{SAT}(\mathcal{KB}) = \text{p-mean}_{i=1}^{|\mathcal{KB}|}(\phi_i) \quad (28)$$

The evaluation context maintains EMA-smoothed satisfaction scores with per-constraint violation tracking for monitoring and diagnostics.

Logical Loss.

$$\mathcal{L}_{\text{logic}} = 1 - \text{SAT}(\mathcal{KB}) \quad (29)$$

As a complement to the LTN constraint system, the Prefrontal module also incorporates quantum-inspired portfolio optimization (Mugel, Lizaso, and Orus, 2022): a QAOA (Quantum Approximate Optimization Algorithm) simulator for combinatorial asset selection via QUBO formulation, VQE (Variational Quantum Eigensolver) for continuous weight optimization using parameterized circuits, and simulated quantum annealing for escaping local minima in non-convex portfolio landscapes. These techniques augment classical Mean-Variance, Risk Parity, and Black-Litterman optimizers, providing additional search capacity in the combinatorially complex space of multi-asset allocation under LTN-enforced constraints.

5.2 Action Selection via Opponent Actor Learning

Action selection in JANUS is designed to be mediated by a model of the Basal Ganglia, comprising Direct (“Go”) and Indirect (“No-Go”) pathways (Collins and Michael J Frank, 2014; Foster, Morris, and Dayan, 2013). This architecture, known as Opponent Actor Learning (OpAL), creates dynamic risk tolerance that adapts to market conditions. Jaskir and Michael J. Frank (2023) extended OpAL with dynamic dopamine modulation, normalized prediction errors, and uncertainty-adjusted learning rates (OpAL*), demonstrating robust advantages in sparse reward environments and large action spaces—conditions directly analogous to financial markets. **No prior work applies OpAL or any basal ganglia computational model to financial decision-making.**

Dopamine modulation governs the balance between pathways: high dopamine (bull market/high confidence) amplifies the Direct pathway, while low dopamine (bear market/uncertainty) amplifies the Indirect pathway. This creates dynamic risk tolerance that adapts to market volatility, implementing allostatic regulation (Sterling, 2012) rather than simple homeostatic feedback. The Hypothalamus module (Section 5.3) translates regime detection into calibrated risk parameters that modulate dopaminergic signaling throughout the decision pipeline.

5.2.1 Dual-Pathway Architecture

Direct Pathway (Go Signal). Encodes the benefits of an action, amplified by dopamine during high-confidence regimes (Schultz, Dayan, and Montague, 1997):

$$Q_{\theta}^{+}(s, a) = \mathbb{E}[\text{Reward} \mid s, a] \quad (30)$$

The implementation includes dopamine sensitivity parameters, learning rates for value updates, decay rates, and threshold-based action release.

Indirect Pathway (No-Go Signal). Encodes the costs and risks, amplified during uncertainty (Collins and Michael J Frank, 2014):

$$Q_{\theta}^{-}(s, a) = \mathbb{E}[\text{Risk} \mid s, a] \quad (31)$$

The implementation provides risk assessment, inhibition generation, and caution scoring.

Action Selection. The final action is determined by the competition between pathways:

$$\mathbf{a}_t = \text{softmax}(\mathbf{d}_{\text{direct}} - \lambda \cdot \mathbf{d}_{\text{indirect}}) \quad (32)$$

where $\lambda > 0$ is the inhibition weight parameter, and each pathway is computed as:

$$\mathbf{d}_{\text{direct}} = \text{ReLU}(\mathbf{W}_{\text{direct}}\mathbf{h} + \mathbf{b}_{\text{direct}}) \quad (33)$$

$$\mathbf{d}_{\text{indirect}} = \text{ReLU}(\mathbf{W}_{\text{indirect}}\mathbf{h} + \mathbf{b}_{\text{indirect}}) \quad (34)$$

where $\mathbf{h}$ is the fused state representation from the Thalamus.

An actor-critic framework ties both pathways together, with decision confidence scoring determining whether the action is released for execution.

5.2.2 Brain Wiring Pipeline

In production, decisions traverse a six-stage pipeline that chains brain regions together:

1. **Regime Detection:** Ensemble classifier identifies the current market state
2. **Hypothalamus:** Adjusts position sizing and risk appetite based on regime and portfolio homeostasis
3. **Amygdala:** Evaluates threat signals; may trigger circuit breakers before further processing
4. **Gating:** Thalamic attention gates filter and weight modality signals
5. **Correlation:** Cross-asset correlation tracking across monitored pairs informs diversification
6. **Execution:** Basal ganglia action selection routes to the Execution Service

This pipeline is operationally richer than the abstract dual-pathway model, reflecting the biological reality that action selection involves coordination across multiple brain regions.

5.3 Homeostatic Risk Regulation

The Hypothalamus implements allostatic regulation (Sterling, 2012), maintaining the system’s internal balance across varying market conditions through adaptive position sizing and risk appetite control:

- **Position Sizing:** Dynamically adjusts position sizes based on regime, portfolio heat, and recent performance using Kelly Criterion-inspired scaling
- **Homeostasis:** Monitors portfolio-level vital signs (exposure, drawdown, correlation) and applies corrective adjustments to maintain target ranges
- **Energy Management:** Tracks “metabolic” state of the trading system—capital utilization, margin usage, and recovery capacity—and modulates aggression accordingly
- **Risk Appetite:** Integrates regime signals from the ensemble detector with internal portfolio state to produce a single risk appetite scalar that modulates all downstream position sizing

The Hypothalamus sits at the second stage of the brain wiring pipeline (Section 5.2), translating raw regime detection into calibrated risk parameters before signals reach the Amygdala and downstream modules.

5.4 Threat Detection and Circuit Breakers

The Amygdala provides rapid threat detection and fear learning (LeDoux, 2000), with connections to substantia nigra enabling fear extinction (Monfils et al., 2009). The amygdala’s role in encoding emotional significance and modulating memory consolidation maps directly to JANUS’s circuit-breaker and fear-network architecture.

5.4.1 Anomaly Detection

Anomaly detection uses multiple complementary methods: Z-score deviation, isolation forest scoring, moving average deviation, percentile outliers, and multivariate scoring. Threats are classified into severity levels (None → Low → Medium → High → Critical).

Mahalanobis Distance.

$$D_M(\mathbf{s}_t) = \sqrt{(\mathbf{s}_t - \boldsymbol{\mu})^T \boldsymbol{\Sigma}^{-1} (\mathbf{s}_t - \boldsymbol{\mu})} \quad (35)$$

where $\boldsymbol{\mu}$ is the historical mean state and $\boldsymbol{\Sigma}$ is the covariance matrix.

Circuit Breaker Condition.

$$\text{Trigger} = \begin{cases} 1 & \text{if } D_M(\mathbf{s}_t) > \tau_{\text{danger}} \\ 0 & \text{otherwise} \end{cases} \quad (36)$$

where τ_{danger} is calibrated to a false-positive rate (e.g., $\tau = 5$ for $p < 0.001$).

Additional Threat Signals.

- Sudden volatility spike: $\sigma_t > 3 \cdot \sigma_{\text{baseline}}$
- Drawdown threshold: cumulative loss $> L_{\text{max}}$
- Liquidity crisis: bid-ask spread $> 10 \times$ normal
- Regime shift detection
- Correlation breakdown across monitored pairs
- Black swan event detection
- Flash crash detection via VPIN

5.4.2 Flow Toxicity and Circuit Breakers

The VPIN (Volume-Synchronized Probability of Informed Trading) calculator (Easley, Lopez de Prado, and O'Hara, 2011; Easley, De Prado, and O'Hara, 2012) uses volume bucket aggregation, bulk volume classification, and rolling window computation to produce a toxicity score. High and critical thresholds trigger graduated responses from position reduction to full kill switch activation.

Production Circuit Breakers.

- **Kill Switch:** Dual-layer design with in-process `AtomicBool` for zero-latency local halt and Redis-backed distributed coordination across services
- **Position Freeze:** Prevents new position entry while allowing exits
- **Safe Mode:** Reduces system to minimal-risk operation
- **Cancel All:** Emergency cancellation of all pending orders

5.4.3 Fear Extinction

The fear learning system implements extinction mechanisms (Monfils et al., 2009; Caligiore et al., 2019), allowing the system to “unlearn” fear when threats have passed. This prevents the agent from becoming permanently paralyzed by a single traumatic market event (e.g., flash crash) while maintaining protective circuit breakers for genuine systemic risks.

5.5 Execution via Cerebellar Forward Models

The Cerebellar module simulates market dynamics to predict execution outcomes before committing to a trade (Wolpert, Miall, and Kawato, 1998; Sokolov, Miall, and Ivry, 2017).

5.5.1 Market Impact Prediction

Following the Almgren-Chriss framework (Almgren and Chriss, 2001; Markwick, 2023), the model predicts slippage and volatility to determine optimal execution trajectories. To detect predatory environments, JANUS employs the VPIN metric (Easley, Lopez de Prado, and O’Hara, 2011; Easley, De Prado, and O’Hara, 2012), which serves as a proxy for flow toxicity—the probability that the counterparty has superior information (see Section 5.4 for the full VPIN description and threat-response architecture). High VPIN levels often precede flash crashes and feed into the Amygdala circuit for threat detection.

$$\hat{p}_{t+1} = f_{\text{cerebellum}}(\mathbf{s}_t, \mathbf{a}_t) \quad (37)$$

5.5.2 Execution Error Correction

The Cerebellum also provides closed-loop error correction for execution quality, implemented through a PID controller:

$$u(t) = K_p \cdot e(t) + K_i \cdot \int_0^t e(\tau) d\tau + K_d \cdot \frac{de(t)}{dt} \quad (38)$$

where $e(t)$ is the deviation between predicted and realized execution cost. Adaptive correction and feedback loops continuously refine the forward model’s predictions based on observed outcomes.

6 Memory and Learning

JANUS’s “sleep state” implements the Complementary Learning Systems (CLS) theory (James L McClelland, Bruce L McNaughton, and O’Reilly, 1995; Kumaran, Hassabis,

and James L. McClelland, 2016), which posits that intelligent agents require two interacting learning systems: a fast-learning hippocampal system for episodic details and a slow-learning neocortical system for statistical generalization. The Backward service runs on a cold path during off-market hours, performing computationally intensive operations to distill daily experiences into long-term knowledge, effectively replicating the biological process of memory consolidation during sleep (Buzsáki, 2015).

6.1 Three-Timescale Memory Hierarchy

The three-tier memory architecture directly implements CLS theory (James L McClelland, Bruce L McNaughton, and O'Reilly, 1995), preventing catastrophic forgetting while enabling rapid learning of new market patterns. Recent work by Masset et al. (2025) provides biological validation for multi-timescale memory architectures, demonstrating that neural systems maintain distinct memory traces across fast, medium, and slow timescales to balance plasticity and stability.

The three tiers mirror this principle. *Short-term* memory (Hippocampal Episodic Buffer) stores raw experiences during trading, capturing full episodic detail without interfering with consolidated knowledge. *Medium-term* consolidation (Sharp-Wave Ripple Replay) replays and compresses salient experiences during post-market “sleep,” prioritized by surprise and salience. *Long-term* memory (Neocortical Schema Formation) distills replayed episodes into stable regime prototypes—abstract schemas that generalize across market conditions while remaining robust to distributional drift.

This separation ensures that JANUS can learn rapidly from novel market events (hippocampal fast-binding) without corrupting the slow-learned statistical regularities that anchor its regime models (neocortical consolidation), a tradeoff that monolithic replay buffers cannot achieve.

6.2 Short-Term: Hippocampal Episodic Buffer

6.2.1 Episodic Buffer

The hippocampal buffer stores episodic memories of individual trading events, enabling fast learning without interfering with consolidated knowledge (James L McClelland, Bruce L McNaughton, and O'Reilly, 1995). Raw experiences are recorded during trading, mirroring the role of the biological hippocampus in episodic memory formation:

$$\mathcal{D}_{\text{hippo}} = \{(s_t, a_t, r_t, s_{t+1}, \mathbf{c}_t, \mathbf{e}_t)\}_{t=1}^T \quad (39)$$

where $\mathbf{c}_t$ contains contextual metadata (volatility, spreads, volume) and $\mathbf{e}_t$ contains emotional tags (fear level, confidence, surprise) that bias consolidation priority.

6.2.2 Pattern Separation and Spatial Mapping

To ensure diverse encoding and avoid interference between similar episodes, the buffer applies pattern separation via random projections:

$$\mathbf{h}_t = \tanh(\mathbf{W}_{\text{rand}} \cdot [s_t; a_t; \mathbf{c}_t]) \quad (40)$$

Experiences are organized into a spatial map that preserves topological relationships between market states, analogous to hippocampal place cells. This enables efficient retrieval of contextually similar experiences during replay.

To address training data scarcity, JANUS augments the episodic buffer with diffusion-based synthetic data generation. A regime-conditional Denoising Diffusion Probabilistic Model (DDPM) (Coletta et al., 2024) generates synthetic market sequences for training data augmentation, with configurable noise schedules (linear, cosine, quadratic), Min-SNR loss weighting, EMA model averaging, and automated quality assessment comparing synthetic vs. real feature distributions and autocorrelation structure.

6.3 Medium-Term: Sharp-Wave Ripple Replay

Consolidation occurs during Sharp-Wave Ripples (SWRs)—high-frequency oscillations that replay compressed sequences of neural activity (Buzsáki, 2015). JANUS mimics this using Prioritized Experience Replay (Schaul et al., 2015), modified to incorporate surprise, emotion, and logical violations. This design is directly inspired by hippocampal sharp-wave ripple replay (M. A. Wilson and Bruce L. McNaughton, 1994; Buzsáki, 2015; Buzsáki, 1989), which the deep RL community adopted as “experience replay” (Mnih et al., 2015; Schaul et al., 2015).

6.3.1 Replay Prioritization

During “sleep” (post-market hours), experiences are replayed in priority order. Biological research shows that replay is biased towards salient and novel events (Joo and L. M. Frank, 2023; Mattar and Nathaniel D. Daw, 2018), which JANUS replicates through composite priority scoring.

Compute TD-error based priority:

$$p_i = |\delta_i| + \epsilon \quad (41)$$

where $\delta_i = r_i + \gamma \max_{a'} Q(s_{i+1}, a') - Q(s_i, a_i)$ and $\epsilon = 10^{-6}$ ensures numerical stability.

Sampling probability:

$$P(i) = \frac{p_i^\alpha}{\sum_j p_j^\alpha} \quad (42)$$

where $\alpha \in [0, 1]$ controls prioritization strength (default $\alpha = 0.6$).

Importance sampling correction:

$$w_i = \left(\frac{1}{N \cdot P(i)} \right)^\beta \quad (43)$$

where β is annealed from $0.4 \rightarrow 1.0$ during training via a configurable increment parameter, fully correcting bias at convergence.

The replay buffer uses a sum tree data structure for $\mathcal{O}(\log n)$ sampling, enabling efficient prioritized replay even with large buffer sizes.

Algorithm 2 Prioritized Experience Sampling

```

1: Input: Buffer  $\mathcal{B} = \{(e_i, p_i)\}_{i=1}^N$ , batch size  $B$ , exponents  $\alpha, \beta$ 
2: Output: Sampled batch with importance weights
3:
4: Compute probabilities:  $P(i) = p_i^\alpha / \sum_j p_j^\alpha$ 
5: for  $k = 1$  to  $B$  do
6:   Sample index  $i_k \sim \text{Categorical}(P)$  via sum-tree traversal
7:   Compute weight:  $w_{i_k} = \left( \frac{1}{N \cdot P(i_k)} \right)^\beta$ 
8:   Add  $(e_{i_k}, w_{i_k})$  to batch
9: end for
10: Normalise weights:  $w_{i_k} \leftarrow w_{i_k} / \max_j w_j$ 
11: return batch with normalised importance weights

```

6.3.2 Sleep-Phase Consolidation

The SWR simulator progresses through biologically inspired phases:

$$\text{Phase} \in \{\text{Awake} \rightarrow \text{Light} \rightarrow \text{Deep} \rightarrow \text{Integration} \rightarrow \text{Transition}\} \quad (44)$$

Each phase adjusts replay parameters (compression ratio, replay rate, consolidation threshold), mirroring the empirical finding that different sleep stages serve distinct memory functions.

6.4 Long-Term: Neocortical Schema Formation

6.4.1 Schema Representation

Schemas represent learned market regime prototypes. The production implementation uses feature-range matching with weighted multi-criteria scoring rather than pure

centroid-based clustering:

$$\text{match}(\mathbf{x}, \mathcal{S}_k) = \frac{1}{|\mathcal{F}|} \sum_{f \in \mathcal{F}} w_f \cdot \mathbb{I}[x_f \in \text{range}_f^{(k)}] \quad (45)$$

where $\mathcal{F}$ is the feature set, w_f is the feature weight, and $\text{range}_f^{(k)}$ is the acceptable range for feature f in schema k . This approach is deterministic and interpretable, providing explicit decision boundaries for each regime.

The system supports seven regime schemas (Bull, Bear, Ranging, Crisis, Recovery, Bubble, Deflation), each with a Markov transition matrix for regime prediction:

$$P(\mathcal{S}_{t+1} = j \mid \mathcal{S}_t = i) = \mathbf{T}_{ij} \quad (46)$$

For embedding-based operations (similarity search, schema formation from new experiences), centroid-based representations remain available:

$$\mathbf{z}_k = \frac{1}{|\mathcal{C}_k|} \sum_{i \in \mathcal{C}_k} \mathbf{h}_i \quad (47)$$

6.4.2 Recall-Gated Consolidation

Only successfully recalled experiences are permitted to update schemas, ensuring that noisy or anomalous episodes do not corrupt long-term knowledge:

$$\mathbf{z}_k \leftarrow \mathbf{z}_k + \eta \cdot \mathbb{I}[\text{recall_success}] \cdot (\mathbf{h}_{\text{new}} - \mathbf{z}_k) \quad (48)$$

Algorithm 3 Schema Formation and Update

- 1: **Input:** Experiences $\mathcal{E} = \{e_1, \dots, e_N\}$, number of clusters K
 - 2: **Output:** Updated schema database
 - 3:
 - 4: Extract embeddings: $\mathbf{h}_i = \text{Embed}(e_i)$ for $i = 1, \dots, N$
 - 5: Cluster: $\mathcal{C} \leftarrow \text{K-means}(\{\mathbf{h}_i\}, K)$ ▷ deep consolidation phase only
 - 6: **for** $k = 1$ to K **do**
 - 7: Compute centroid: $\mathbf{z}_k = \frac{1}{|\mathcal{C}_k|} \sum_{i \in \mathcal{C}_k} \mathbf{h}_i$
 - 8: Compute statistics: $n_k = |\mathcal{C}_k|$, $\bar{r}_k = \frac{1}{|\mathcal{C}_k|} \sum_{i \in \mathcal{C}_k} r_i$
 - 9: **Upsert** schema k in Qdrant with vector $\mathbf{z}_k$ and metadata $(n_k, \bar{r}_k)$
 - 10: **end for**
 - 11: **Note:** Online recall-gated updates use the EMA rule above; K-means runs only during deep sleep.
-

6.5 UMAP Visualization and Drift Detection

To visualize and manage schema structures, JANUS employs Uniform Manifold Approximation and Projection (UMAP) (McInnes, Healy, and Melville, 2018b), which preserves both local and global structure better than alternatives like t-SNE.

AlignedUMAP (McInnes, Healy, and Melville, 2018b; McInnes, Healy, and Melville, 2018a) tracks how internal representations evolve over time by aligning manifolds across different time steps to monitor representational drift, maintaining consistent embeddings across sleep cycles. The full UMAP loss includes both attraction and repulsion terms:

$$\mathcal{L}_{\text{UMAP}} = \sum_{i \neq j} [w_{ij} \log(q_{ij}) + (1 - w_{ij}) \log(1 - q_{ij})] \quad (49)$$

where $q_{ij} = (1 + \|\mathbf{y}_i - \mathbf{y}_j\|^2)^{-1}$.

Note: In practice, the repulsion term $(1 - w_{ij})$ is approximated via *negative sampling* to achieve $\mathcal{O}(N)$ complexity. For each positive edge, we sample $k = 5$ random negative pairs.

For real-time monitoring, a parametric UMAP model trains a neural network to project new experiences:

$$\mathbf{y}_{\text{new}} = f_{\theta}(\mathbf{h}_{\text{new}}) \quad (50)$$

The parametric UMAP implementation uses a multi-layer encoder network trained on the fuzzy simplicial set graph, with configurable distance metrics (Euclidean, Cosine, Manhattan), smooth k -nearest-neighbor bandwidth estimation, and a Qdrant bridge for persistent storage of projected embeddings. A drift detection module computes embedding-space displacement between reference and live data, classifying severity (Low/Moderate/High/Critical) for regime monitoring. Regime cluster analysis provides per-cluster statistics (centroid, intra-cluster distances, standard deviation) and inter-cluster separation metrics.

7 Implementation

JANUS is implemented as a pure Rust stack, from training and inference to production services and deployment orchestration. This section covers implementation-specific details—service topology, execution infrastructure, strategy configuration, and operational safety—that are not addressed in the preceding mathematical specifications (§4–6).

7.1 Rust-Native ML Stack

The exclusive use of Rust across the entire stack—training, inference, and production services—is justified by the language’s formal safety guarantees (Matsakis and Klock, 2014; Jung et al., 2018) and the performance requirements of high-frequency trading systems:

1. **Performance:** Zero-cost abstractions, no GC pauses—critical for sub-microsecond latency
2. **Safety:** Memory safety without runtime overhead, preventing undefined behavior
3. **Concurrency:** Fearless async/await with Tokio (Tokio Contributors, 2024) for handling high-frequency websocket feeds
4. **Ecosystem:** End-to-end ML training and inference via Candle (Hugging Face, 2024)
5. **Unified Stack:** Single language for the entire pipeline eliminates cross-language serialization overhead, deployment complexity, and FFI boundary risks

The following Rust-native frameworks comprise the JANUS ML stack:

Framework	Pros	Cons	Use Case
Candle (Hugging Face, 2024)	Pure Rust, HuggingFace integration, minimal deps, autograd	Younger ecosystem	Primary training & inference (DiffGAF, ViViT, LTN, DQN)
ndarray	Zero-dependency numerical arrays, mature ecosystem	No autograd	Traditional fuzzy logic, GAF encoding, feature engineering
Polars (Polars Contributors, 2024)	High-speed DataFrames	N/A	Data manipulation

The system operates as a **fully Rust-native** ML pipeline with no external language dependencies:

- End-to-end training and inference in Rust via Candle with autograd support
- Custom differentiable kernels for DiffGAF and LTN operations
- Double DQN with online/target networks for reinforcement learning

- Native model serialization using `safetensors` format
- GPU acceleration via `wgpu` (Candle) and CUDA (optional)
- Zero cross-language overhead—no FFI bridges, no serialization boundaries
- Training infrastructure: AdamW/SGD optimizers, warmup+cosine LR scheduling, prioritized replay buffers

Fractal Signal Processing (DSP)

A dedicated `dsp` crate provides a sub-microsecond signal processing pipeline that runs ahead of the main neural pipeline and feeds regime detection. The implementation applies the Fractal Adaptive Moving Average (FRAMA) together with the Sevcik fractal dimension estimator to produce an 8-dimensional feature vector per tick: divergence, alpha, fractal dimension, Hurst exponent, regime, sign, deviation, and confidence. At sub-microsecond latency per tick, the DSP pipeline enables high-frequency regime awareness without adding latency to the hot path. All computations use incremental $\mathcal{O}(1)$ updates; no look-back window is stored beyond the minimum required for each indicator.

The `indicators` crate complements the DSP pipeline with classical technical indicators (ADX, ATR, Bollinger Bands, EMA, RSI, MACD), all implemented with the same incremental $\mathcal{O}(1)$ per-tick computation guarantee. The `compliance` crate provides the `WashSaleDetector` (full 30-day lookback/lookforward, partial wash sale handling, cost basis tracking) and `ComplianceSheriff` (daily loss limits, maximum loss thresholds, mandatory stop-losses), both operating independently of the LTN constraint layer to provide hard regulatory floors.

7.2 Service Architecture and Deployment

7.2.1 Service Topology

The system comprises eight services plus supporting infrastructure (see Figure in Section 3.4 for the component diagram). In the current production deployment these run as modules of a single unified binary: a supervisor process spawns the Data, Forward, Backward, CNS, and API modules in-process (with per-module health reporting and restart policies), which eliminates inter-service serialization on the hot path while preserving the logical boundaries below. The standalone-service topology remains supported for horizontally scaled deployments.

1. **Forward Service:** Ports gRPC 50051, HTTP 7000. Dependencies: native model artifacts, Redis, PostgreSQL. Resource limits: 2GB memory, 2 CPU cores.

2. **Backward Service:** Internal service (no external ports). Dependencies: PostgreSQL, Redis. Scheduling: triggered during market close.
3. **Execution Service:** Ports HTTP 8081, gRPC 50052. Dependencies: exchange API credentials, Redis (kill switch). Multi-exchange order routing.
4. **Data Service:** Internal service for centralized market data management. Dependencies: exchange WebSocket feeds, QuestDB (time series).
5. **CNS Service:** Internal service for health monitoring and preflight validation. Dependencies: all other services (monitoring target).
6. **API Service:** HTTP API gateway for external access. Dependencies: Forward Service, Backward Service.
7. **Registry Service:** Internal service discovery and registration. Dependencies: Redis.
8. **Optimizer Service:** Hyperparameter optimization with backtesting integration (metrics on the unified :9090 endpoint in single-binary mode).

Infrastructure:

- **PostgreSQL** (port 5432): Primary relational storage for signals, portfolios, and performance
- **Redis** (port 6379): Operational state, kill switch coordination, hot-reloadable config
- **QuestDB** (ports 9000/9009): High-performance time series storage for market data
- **Qdrant:** Vector similarity search engine for schema pattern matching
- **Prometheus** (port 9090): Metrics collection with a curated alert-rule set, audited so that every rule references an exported metric (data-quality SLOs, host safety, per-bot performance)
- **Grafana** (port 3000): Visualization dashboards (signal pipeline, data-quality SLOs, operations, per-bot performance)
- **Alertmanager** (port 9093): Alert routing with Discord integration
- **Jaeger** (port 16686): Distributed tracing
- **Loki + Promtail:** Centralized log aggregation with label-based organization

- **Nginx:** Reverse proxy with TLS termination and internal-token injection (deployments sit behind a WireGuard-based private network)

Service Communication:

$$\text{Forward} \xrightarrow[\text{gRPC}]{\text{signals}} \text{Execution} \xrightarrow[\text{HTTP}]{\text{orders}} \text{Exchanges} \quad (51)$$

$$\text{Forward} \xrightarrow[\text{SQL}]{\text{experiences}} \text{PostgreSQL} \xrightarrow[\text{nightly}]{\text{batch}} \text{Backward} \xrightarrow[\text{SQL}]{\text{schemas}} \text{PostgreSQL} \quad (52)$$

Volume Management: Model artifacts are stored on shared read-only volumes; PostgreSQL and QuestDB use persistent volumes with automated backups; Redis data is ephemeral with configurable persistence; experience buffers rotate daily.

7.2.2 Execution Service

The Execution Service is a dedicated microservice responsible for multi-exchange order routing, providing a clean separation between decision-making (Forward Service) and order management. It supports multiple cryptocurrency exchanges through a unified adapter interface:

- **Kraken:** Full REST and WebSocket adapter with order normalization (primary exchange)
- **Bybit:** Dedicated client crate with unified order types
- **Coinbase:** REST adapter with authentication and rate limiting
- **OKX:** REST adapter with order format translation
- **Binance:** Legacy connector maintained for backward compatibility
- **Kucoin:** Legacy connector maintained for backward compatibility

Each adapter implements a common `ExchangeAdapter` trait, enabling transparent exchange selection at runtime. The service provides three production execution algorithms: TWAP (Time-Weighted Average Price with configurable slice intervals), VWAP (Volume-Weighted Average Price with volume profile estimation), and Iceberg (hidden large orders exposing small visible tip orders to minimize market impact), complemented by per-venue latency tracking, fill rate monitoring, and slippage analysis.

7.3 Trading Strategies and Regime-Aware Gating

The `strategies` crate implements nine regime-aware trading strategies, each designed for specific market conditions as identified by the ensemble regime detector.

Trend-Following Strategies:

- **EMA Flip:** 8/21 EMA crossover with ATR-based stops, trading pullbacks to the fast EMA in the trend direction
- **EMA Ribbon Scalper:** 8/13/21 EMA ribbon with volume confirmation for higher-quality pullback entries
- **Trend Pullback:** Fibonacci retracement entries within established trends, confirmed by RSI divergence and candlestick patterns (pin bars, engulfing)
- **Momentum Surge:** Detects sudden price surges with volume spikes, entering on the first pullback within the surge
- **Multi-Timeframe Trend:** EMA 50/200 crossover with ADX strength and higher-timeframe alignment

Mean-Reversion Strategies:

- **Mean Reversion:** Bollinger Bands with RSI confirmation and ATR-based stops
- **Bollinger Squeeze Breakout:** Detects low-volatility squeeze periods and generates breakout signals when price escapes the bands
- **VWAP Scalper:** Mean reversion scalping around the Volume-Weighted Average Price with standard deviation bands
- **Opening Range Breakout:** Trades breakouts above or below the first N candles of a session with volume confirmation

Strategy Gating and Affinity. The `StrategyGate` module controls which strategies execute based on:

- **Regime Compatibility:** Each asset can define preferred strategies per regime (e.g., Trending → EMA Flip, MeanReverting → Bollinger Squeeze)
- **Affinity Scoring:** Strategy-asset affinity tracker with performance-weighted scoring
- **Allowlists/Denylists:** Per-asset strategy filtering via TOML configuration
- **Untested Strategy Policy:** Configurable flag to allow or block strategies without historical performance data

7.4 Safety Infrastructure

The Central Nervous System (CNS) Service provides system-wide health monitoring and operational safety, analogous to the autonomic nervous system's regulation of vital functions.

7.4.1 Preflight Validation

Before entering live or paper trading, the CNS Service executes a five-phase preflight sequence, each with configurable criticality levels (Critical, Required, Optional):

1. **Infrastructure Phase:** Database connectivity, Redis availability, Qdrant health, shared memory paths
2. **Sensory Phase:** Exchange API authentication, WebSocket feed health, market data freshness
3. **Regulatory Phase:** Kill switch state verification, compliance rule loading, wash sale detector initialization
4. **Strategy Phase:** Model file availability, version consistency, strategy affinity configuration
5. **Executive Phase:** Memory and CPU resource adequacy, neuromorphic brain region initialization

Critical failures abort the boot sequence; required failures block trading but allow monitoring; optional failures are logged but permit full operation. The `PreFlightRunner` supports both sequential and parallel-within-phase execution modes, producing a comprehensive `BootReport` with Discord-formatted notifications.

7.4.2 Runtime Monitoring

During operation, the CNS Service provides:

- **Watchdog Monitoring:** Heartbeat-based component liveness detection with configurable degraded/dead thresholds, per-component criticality levels (Critical, Important, NonEssential), and automatic kill switch triggering when critical components die
- **Boot Reports:** Comprehensive system state summary with pass/fail/skip counts per phase
- **Prometheus Metrics:** 15+ custom metrics exposed for Grafana dashboards (6+ pre-built dashboards)
- **Alert Integration:** Slack, PagerDuty, and generic webhook notifications with severity-based routing
- **Distributed Tracing:** Jaeger integration for cross-service latency analysis

- **Circuit Breakers:** Per-component circuit breakers with Closed → Open → Half-Open state machine, configurable failure windows, and recovery timeouts
- **Reflex Actions:** Automated responses including component restart, throttling, graceful shutdown, and safe command execution with allowlist validation
- **Neuromorphic Brain Coordinator:** Topological initialization ordering of all 10 brain regions based on dependency graphs, per-region health scoring, and global brain activation/deactivation

8 Limitations and Open Problems

No system of JANUS's complexity can be presented without candid engagement with its limitations.

8.1 Research-to-Production Maturity Gap

The subsystems described in this paper differ in deployment maturity, and we state the gap plainly. The live trading path is a deterministic regime/rule engine—technical indicators → regime detection → homeostatic position scaling → threat filtering → execution gating—and the brain wiring pipeline (Section 5.2) that orchestrates it is rule-based rather than neural. Its safety, compliance, and signal-processing infrastructure (kill switch, wash-sale/prop-firm compliance, VPIN flow-toxicity, the fractal DSP pipeline, ensemble regime detection, and Wilson-Cowan thalamic dynamics) is implemented and live, as is its observability layer: generated signals are persisted from the in-process signal bus to Redis for the read APIs and operator UI, and a data-completeness SLI computed in the live ingestion path backs a data-loss SLO alert (which, on its first armed day, surfaced a silently dead exchange stream caused by an upstream ticker migration). By contrast, the neuromorphic and deep-learning subsystems that constitute the cognitive core—OpAL action selection (Section 5.2), the LTN constraint engine (Section 5.1), the DiffGAF/ViViT vision stack (Section 4), and the three-timescale memory hierarchy (Section 6)—are implemented as standalone modules and validated in isolation, but are not yet integrated into the live inference loop, and no trained model weights currently drive live decisions. Wiring these modules into the hot path, exporting trained weights, and validating end-to-end behavior is the primary thrust of ongoing work; Section 10.3 gives a per-component breakdown of maturity.

8.2 The Brain Metaphor Critique

Every generation projects its latest technology onto the brain: clocks, telephone switchboards, computers, and now neural networks. Critics may argue that JANUS's brain-

region mapping is the latest instance of this pattern. We acknowledge the validity of this concern and offer two defenses:

1. **Functional, not literal:** Following Hassabis et al. (2017), JANUS uses neuroscience as *functional inspiration*, not biological simulation. The mappings are justified by computational principles (e.g., reinforcement learning for action selection, supervised error correction for timing) rather than cellular fidelity.
2. **Scientifically productive:** The Doya framework (Doya, 1999) has generated falsifiable predictions confirmed by neuroimaging. The super-learning hypothesis (Caligiore et al., 2019) demonstrates that multi-region integration produces capabilities exceeding single-region models. JANUS’s architecture is grounded in this validated framework.

8.3 Integration Complexity

Glasmachers (2017) demonstrated experimentally that end-to-end training can fail even for small multi-module systems due to non-trivial couplings between components. JANUS integrates ten brain-region modules, creating substantial coupling risk.

Mitigations: (i) Modular training: each brain region can be trained independently before integration. (ii) Service boundaries: the Forward and Backward services communicate via gRPC, creating natural decoupling points. (iii) The brain wiring pipeline (Section 5.2) defines explicit information flow, preventing uncontrolled coupling.

Honest assessment: Integration testing at scale remains an open challenge. Ablation studies (Section 9) are needed to quantify inter-region dependencies.

8.4 Non-Stationarity and Regime Shifts

Financial environments are fundamentally non-stationary. RL algorithms optimized for one regime may fail catastrophically in the next (Padakandla et al., 2020). JANUS’s ensemble regime detector is an important mitigation but is itself a heuristic—novel regimes not represented in training data may not be correctly classified.

Mitigations: (i) Three-timescale memory allows rapid adaptation at short timescales while preserving long-term knowledge (Masset et al., 2025). (ii) Brain-inspired replay (Ven, Siegelmann, and Tolias, 2020) helps mitigate catastrophic forgetting. (iii) The hypothalamic homeostatic regulator provides regime-independent risk controls.

Honest assessment: No existing approach fully solves the stability–plasticity dilemma in non-stationary environments. JANUS’s multi-timescale architecture is theoretically motivated but requires empirical validation across diverse market conditions.

8.5 Backtest Overfitting Risk

Bailey et al. (2015) established the Combinatorially Symmetric Cross-Validation (CSCV) framework showing that conventional backtesting dramatically overstates expected performance. López de Prado (2018b) warned that “ML algorithms will always find a pattern, even if there is none.” With ten brain regions and numerous hyperparameters, JANUS is particularly susceptible to overfitting.

Mitigations: (i) Walk-forward validation protocol described in Section 9. (ii) LTN constraints provide domain knowledge that regularizes against spurious patterns. (iii) The CNS service monitors for performance degradation relative to out-of-sample benchmarks.

8.6 Crowding Resistance Is Not Crowding Immunity

As discussed in Section 3.5, the literature supports heterogeneity as a mechanism for crowding *resistance*, not crowding *immunity*. Key limitations:

- Khandani and Lo (2011) showed that independently developed strategies converged on similar factors in 2007 despite nominal diversity
- Evolutionary dynamics can cause heterogeneous populations to converge over time if selection pressure favors similar strategies
- Tail events may produce correlated liquidation regardless of strategy diversity (Caccioli et al., 2014)

Mitigations: (i) Active monitoring of cross-instance correlation. (ii) Periodic re-randomization of heterogeneity parameters. (iii) Multi-agent LOB simulation for ongoing validation (Section 9).

8.7 LTN Scalability

Wan et al. (2024) found that LTN workloads exhibit a fundamental hardware utilization challenge: symbolic operations are memory-bounded while neural operations are compute-bounded, creating inefficiencies on standard hardware. As the axiom count grows, LTN inference latency may conflict with hot-path requirements.

Mitigations: JANUS adopts a dual-mode LTN design: strict mode (full axiom evaluation) for the cold path and approximate mode (reduced axiom set with cached evaluations) for the hot path.

8.8 Adversarial Robustness

Goldblum et al. (2021) demonstrated that adversarial traders can fool automated ML systems by placing adversarial orders on public exchanges. JANUS’s neural components—particularly the visual pattern recognition and sentiment analysis modules—are vulnerable to such attacks.

Mitigations: (i) The amygdala’s multi-method anomaly detection provides some defense against adversarial inputs. (ii) The FNI-RL fear network can learn to recognize and avoid adversarial patterns. (iii) LTN constraints provide a logic-based floor that adversarial perturbations cannot violate.

Honest assessment: Adversarial robustness in financial ML is an open research problem. JANUS’s defenses are heuristic, not provably robust.

9 Validation Framework

Empirical claims require a rigorous experimental protocol. This section describes the planned validation framework.

9.1 Multi-Agent LOB Simulation

The primary validation of JANUS’s crowding-resistance claim requires multi-agent simulation using the Rust-native limit order book simulator (Fu, Pakkanen, and Cont, 2024) (CPU matching engine; GPU acceleration is planned).

Experimental conditions:

- **Condition A (Homogeneous):** 50–200 identical JANUS instances with shared parameters trading simultaneously
- **Condition B (Heterogeneous):** 50–200 JANUS instances with engineered heterogeneity (Table 2)
- **Control:** Background liquidity provided by zero-intelligence agents (Farmer, Patelli, and Zovko, 2005)

Metrics:

- Return kurtosis (tail risk under crowding)
- Net signed volume vs. price move (co-impact measurement)
- Preservation of the square-root impact law (Tóth et al., 2011)
- Cross-instance return correlation
- Effective market capacity (Sharpe ratio degradation as N increases)

9.2 Walk-Forward Backtesting

Following Bailey et al. (2015) and López de Prado (2018a), the backtesting protocol uses:

1. **Combinatorially Symmetric Cross-Validation (CSCV):** Partition historical data into S subsets, form all $\binom{S}{S/2}$ train/test splits, and compute the probability of back-test overfitting
2. **Out-of-sample Sharpe ratio degradation:** Compare in-sample to out-of-sample Sharpe ratios across multiple walk-forward windows
3. **Regime-conditional performance:** Disaggregate performance by detected regime to identify regime-specific overfitting

9.3 Ablation Studies

Ablation studies isolate the contribution of each component:

- **Brain region ablation:** Remove each brain region individually and measure performance degradation, directly demonstrating the value of multi-region integration
- **Memory tier ablation:** Compare the three-tier CLS memory against single-tier experience replay
- **LTN ablation:** Remove symbolic constraints and measure constraint violation rates and risk-adjusted returns
- **Heterogeneity ablation:** Compare homogeneous vs. heterogeneous multi-agent performance (overlaps with LOB simulation)
- **Regime detection ablation:** Remove the ensemble regime detector and measure adaptation speed to market transitions

9.4 Benchmark Comparisons

- **Standard RL baselines:** PPO, SAC, and TD3 on the same market data and reward function
- **Rule-based systems:** Classic momentum, mean-reversion, and factor-based strategies (Quant 1.0/2.0 comparisons)
- **Single-region systems:** Pure basal ganglia RL without cerebellar timing, pure replay without CLS consolidation
- **Existing neuromorphic approaches:** SNN-based portfolio optimization (Mohan et al., 2025) where feasible

10 Conclusion

Project JANUS represents a paradigm shift in quantitative trading: from opaque black boxes to transparent, brain-inspired systems that combine the best of deep learning and symbolic reasoning. Grounded in the Doya–Hassabis framework (Doya, 1999; Hassabis et al., 2017) and motivated by the growing threat of strategy crowding (Bucci, Mastromatteo, et al., 2020; Khandani and Lo, 2011), JANUS is—to our knowledge—the first trading system to map multiple brain regions to distinct functional subsystems within a unified architecture.

10.1 Key Innovations

1. **Neuromorphic Architecture:** Ten functionally motivated brain regions with modular design, distributed training infrastructure, and Wilson-Cowan oscillatory dynamics (Doya, 1999; Caligiore et al., 2019)
2. **Crowding Resistance:** Engineered heterogeneity across deployments reducing co-impact costs and cross-instance correlation (Wagner, 2011; DeMiguel, Martin-Utrera, and Uppal, 2021)
3. **Neuro-Symbolic Fusion:** LTN constraint engine with domain-specific axioms and linguistic hedges bridging neural networks and logical constraints (Badreddine et al., 2022)
4. **Multi-Timescale Memory:** Three-tier hierarchy with sleep-phase consolidation mirrors hippocampal–neocortical transfer, implementing CLS theory (Kumaran, Hassabis, and James L. McClelland, 2016; Masset et al., 2025)
5. **OpAL Decision Engine (design):** An Opponent Actor Learning module for financial action selection (Collins and Michael J Frank, 2014; Jaskir and Michael J. Frank, 2023)—to our knowledge the first application of a basal ganglia computational model to this problem
6. **Ensemble Regime Detection:** HMM, statistical, and technical methods fused for robust market state identification with strategy routing
7. **Production-Ready Safety:** Dual-layer kill switch (in-process + Redis) with four breaker actions (freeze, safe-mode, cancel-all), multi-method anomaly detection, FNI-RL fear network, and CNS health monitoring with five-phase preflight
8. **Pure Rust Stack:** End-to-end training and inference in Rust (Matsakis and Klock, 2014) with high-performance, safe, and maintainable eight-service architecture

9. **Fractal Signal Processing:** Sub-microsecond DSP pipeline producing 8D feature vectors via FRAMA and Sevcik fractal dimension for regime-aware signal generation
10. **GPU Compute Infrastructure:** WGSL compute kernels (MatMul, Softmax, LayerNorm, Attention, GELU, Reduce) with a Candle tensor bridge—authored with CPU-reference execution; GPU dispatch pending

10.2 Future Work

The following items remain as planned enhancements beyond the current implementation:

- Hardware acceleration using FPGAs (Marino et al., 2023; Vemeko, 2023; AMD, 2023) and neuromorphic chips for nanosecond-level latency in high-frequency trading applications
- Execution of the full validation framework described in Section 9
- Investigation of adversarial robustness guarantees beyond heuristic defenses

10.3 Implementation Status

Implementation maturity varies by subsystem, and we summarize it here rather than presenting every component as live. The production trading path is a deterministic regime/rule engine (technical indicators → regime detection → homeostatic position scaling → threat filtering → execution gating); its safety, compliance, and signal-processing infrastructure is implemented and live. The neuromorphic and deep-learning subsystems that constitute the cognitive core (Sections 4–6)—notably OpAL action selection, the LTN constraint engine, and the DiffGAF/ViViT vision stack—are implemented as standalone modules and validated in isolation, but are not yet wired into the live inference loop; integrating them is ongoing work (Section 8.1). The table below describes each component at its current maturity, annotated as live, standalone, or scaffolded/planned where the distinction matters.

Component	Description
Flow Toxicity	VPIN calculator with flash crash detection
Risk Control	Full PID controller with Ziegler-Nichols auto-tuning
Compliance	Wash sale detection, proprietary firm rule enforcement
Strategies	Nine regime-aware strategies with gating and affinity scoring
Safety	Kill switch (in-process + Redis) with four breaker actions—freeze/safe-mode/cancel-all; FNI-RL fear network
Training	Multi-GPU/multi-node with AllReduce, Parameter Server, Ring AllReduce
Feudal RL	Cortex manager and hippocampus worker agents
LTN	10 signal-logic axioms with fuzzy-logic loss (standalone)
Regime	Markov transition matrix with stationary distribution
DSP	8D feature vector pipeline (FRAMA + Sevcik), incremental per-tick
Vision	DiffGAF → ViViT pipeline (dual-GAF)—standalone, not yet in live loop
CLS Memory	Three-tier hierarchy with sleep-phase consolidation
Thalamus	Wilson-Cowan oscillation models (H. R. Wilson and Cowan, 1972; Y. Zhang et al., 2024) with Hilbert-transform estimation
Forecasting	Chronos tokenizer (Ansari et al., 2024b; Ansari et al., 2024a) with statistical-baseline forecasting (ONNX serving scaffolded, not wired)
NLP	BERT/FinBERT sentiment via Candle (Hugging Face, 2024) with Qdrant storage
Visualization	Parametric UMAP (McInnes, Healy, and Melville, 2018b) with drift detection
Quantum	QAOA, VQE, simulated annealing (Mugel, Lizaso, and Orus, 2022)
Synthetic Data	Regime-conditional DDPM diffusion models (Coletta et al., 2024) (Min-SNR/EMA scaffolded)
LOB Simulator	CPU matching engine (Fu, Pakkanen, and Cont, 2024) with Almgren-Chriss impact (GPU acceleration planned)
GPU Kernels	WGSL compute kernels (MatMul, Softmax, LayerNorm, Attention, GELU, Reduce)—authored; CPU-reference execution, GPU dispatch pending
Qdrant Client	Circuit breaker, exponential backoff, TLS, mock fallback

Repository & Contact

GitHub: <https://github.com/nuniesmith/janus>

For implementation code, updates, and discussions, visit the repository.

“The god of beginnings and transitions, looking simultaneously to the future and the past.”

References

- Alayrac, Jean-Baptiste, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, Arthur Mensch, Katherine Millican, Malcolm Reynolds, et al. (2022). “Flamingo: A visual language model for few-shot learning”. In: *Advances in Neural Information Processing Systems*. Vol. 35. Introduces gated cross-attention for multimodal fusion.
- Almgren, Robert and Neil Chriss (2001). “Optimal execution of portfolio transactions”. In: *Journal of Risk* 3, pp. 5–40.
- AMD (2023). *AMD Alveo U55C data center accelerator card*. URL: <https://www.amd.com/en/products/accelerators/alveo/u55c.html>.
- Ansari, Abdul Fatir et al. (2024a). “Introducing Chronos-2: Large language model time series forecasting”. In: *Amazon Science Blog*. URL: <https://www.amazon.science/blog/introducing-chronos-2>.
- Ansari, Abdul Fatir, Lorenzo Stella, Caner Turkmen, Xiyuan Zhang, Pedro Mercado, Huibin Shen, Oleksandr Shchur, Syama Syndar Rangapuram, Sebastian Pineda Arango, Shubham Kapoor, et al. (2024b). “Chronos: Learning the language of time series”. In: *arXiv preprint arXiv:2403.07815*.
- Arnab, Anurag, Mostafa Dehghani, Georg Heigold, Chen Sun, Mario Lučić, and Cordelia Schmid (2021). “ViViT: A video vision transformer”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 6836–6846.
- Badreddine, Samy, Artur d’Avila Garcez, Luciano Serafini, and Michael Spranger (2022). “Logic tensor networks”. In: *Artificial Intelligence* 303, p. 103649. DOI: [10.1016/j.artint.2021.103649](https://doi.org/10.1016/j.artint.2021.103649).
- Bailey, David H., Jonathan M. Borwein, Marcos López de Prado, and Qiji Jim Zhu (2015). “The probability of backtest overfitting”. In: *Journal of Computational Finance*. SSRN: 2326253.
- Baltas, Nick (2019). “The impact of crowding in alternative risk premia investing”. In: *Financial Analysts Journal* 75.3, pp. 50–65. DOI: [10.1080/0015198X.2019.1600955](https://doi.org/10.1080/0015198X.2019.1600955).
- Bi, Yingjia and Vince D. Calhoun (2025). “The financial connectome: A brain-inspired framework for modeling latent market dynamics”. In: *arXiv preprint arXiv:2508.02012*.
- Bouchaud, Jean-Philippe, Julius Bonart, Jonathan Donier, and Martin Gould (2018). *Trades, Quotes and Prices: Financial Markets Under the Microscope*. Cambridge University Press.
- Brown, Gregory, Philip Howard, and Christian Lundblad (2022). “Crowded trades and tail risk”. In: *Review of Financial Studies* 35, pp. 3231–3271. DOI: [10.1093/rfs/hh-ab107](https://doi.org/10.1093/rfs/hh-ab107).

- Bucci, Frédéric, Michael Benzaquen, Fabrizio Lillo, and Jean-Philippe Bouchaud (2019). “Crossover from linear to square-root market impact”. In: *Physical Review Letters* 122.10. arXiv: 1811.05230, p. 108302.
- Bucci, Frédéric, Iacopo Mastromatteo, Zoltán Eisler, Fabrizio Lillo, Jean-Philippe Bouchaud, and Charles-Albert Lehalle (2020). “Co-impact: Crowding effects in institutional trading activity”. In: *Quantitative Finance* 20.2, pp. 193–205. DOI: [10.1080/14697688.2019.1660398](https://doi.org/10.1080/14697688.2019.1660398).
- Buzsáki, György (1989). “Two-stage model of memory trace formation: A role for “noisy” brain states”. In: *Neuroscience* 31.3, pp. 551–570. DOI: [10.1016/0306-4522\(89\)90423-5](https://doi.org/10.1016/0306-4522(89)90423-5).
- (2015). “Hippocampal sharp wave-ripple: A cognitive biomarker for episodic memory and planning”. In: *Hippocampus* 25.10, pp. 1073–1188. DOI: [10.1002/hipo.22488](https://doi.org/10.1002/hipo.22488).
- Caccioli, Fabio, Munik Shrestha, Cristopher Moore, and J. Doyne Farmer (2014). “Stability analysis of financial contagion due to overlapping portfolios”. In: *Journal of Banking & Finance* 46, pp. 233–245.
- Caligiore, Daniele et al. (2017). “Consensus paper: Towards a systems-level view of cerebellar function: The interplay between cerebellum, basal ganglia, and cortex”. In: *The Cerebellum* 16, pp. 203–229.
- Caligiore, Daniele, Michael A. Arbib, R. Chris Miall, and Gianluca Baldassarre (2019). “The super-learning hypothesis: Integrating learning processes across cortex, cerebellum and basal ganglia”. In: *Neuroscience & Biobehavioral Reviews* 100, pp. 19–34.
- Chen, Jun-Hao and Yun-Cheng Tsai (2020). “Encoding candlesticks as images for pattern classification using convolutional neural networks”. In: *Financial Innovation* 6.1, p. 26. DOI: [10.1186/s40854-020-00187-0](https://doi.org/10.1186/s40854-020-00187-0).
- Coletta, Andrea et al. (2024). “Conditional generative models for simulation of EMG during naturalistic movements”. In: *arXiv preprint*. Generative diffusion models applied to sequential financial data simulation.
- Collins, Anne GE and Michael J Frank (2014). “Opponent actor learning (OpAL): Modeling interactive effects of striatal dopamine on reinforcement learning and choice incentive”. In: *Psychological Review* 121.3, p. 337. DOI: [10.1037/a0037015](https://doi.org/10.1037/a0037015).
- Daw, Nathaniel D, John P O’Doherty, Peter Dayan, Ben Seymour, and Raymond J Dolan (2006). “Cortical substrates for exploratory decisions in humans”. In: *Nature* 441.7095, pp. 876–879.
- DeMiguel, Victor, Alberto Martin-Utrera, and Raman Uppal (2021). “What alleviates crowding in factor investing?” In: *CEPR Discussion Paper 16527*. SSRN: 3928838.
- Dosovitskiy, Alexey, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, Matthias Minderer, Georg Heigold,

- Sylvain Gelly, et al. (2020). “An image is worth 16x16 words: Transformers for image recognition at scale”. In: *arXiv preprint arXiv:2010.11929*.
- Doya, Kenji (1999). “What are the computations of the cerebellum, the basal ganglia and the cerebral cortex?” In: *Neural Networks* 12.7–8, pp. 961–974. DOI: [10.1016/S0893-6080\(99\)00046-5](https://doi.org/10.1016/S0893-6080(99)00046-5).
- (2002). “Metalearning and neuromodulation”. In: *Neural Networks* 15.4–6, pp. 495–506. DOI: [10.1016/S0893-6080\(02\)00044-8](https://doi.org/10.1016/S0893-6080(02)00044-8).
- Easley, David, Marcos Lopez De Prado, and Maureen O’Hara (2012). “Flow toxicity and liquidity in a high-frequency world”. In: *The Review of Financial Studies* 25.5, pp. 1457–1493.
- Easley, David, Marcos M Lopez de Prado, and Maureen O’Hara (2011). “The microstructure of the “flash crash”: flow toxicity, liquidity crashes, and the probability of informed trading”. In: *The Journal of Portfolio Management* 37.2, pp. 118–128.
- Evans, Jonathan St B. T. and Keith E. Stanovich (2013). “Dual-process theories of higher cognition: Advancing the debate”. In: *Perspectives on Psychological Science* 8.3, pp. 223–241.
- Evans, Jonathan St BT (2008). “Dual-processing accounts of reasoning, judgment, and social cognition”. In: *Annual Review of Psychology* 59, pp. 255–278.
- Ezinwoke, B. and O. Rhodes (2025). “Predicting price movements in high-frequency financial data with spiking neural networks”. In: *arXiv preprint arXiv:2512.05868*.
- Fama, Eugene F and Kenneth R French (1993). “Common risk factors in the returns on stocks and bonds”. In: *Journal of Financial Economics* 33.1, pp. 3–56.
- Farmer, J. Doyne, Paolo Patelli, and Ilija I. Zovko (2005). “The predictive power of zero intelligence in financial markets”. In: *Proceedings of the National Academy of Sciences* 102.6, pp. 2254–2259.
- Foster, David J, Richard GM Morris, and Peter Dayan (2013). “Mechanisms of hierarchical reinforcement learning in corticostriatal circuits”. In: *Nature Neuroscience* 16, pp. 383–391.
- Frank, Michael J, Bryan Loughry, and Randall C O’Reilly (2006). “Making working memory work: a computational model of learning in the prefrontal cortex and basal ganglia”. In: *Neural Computation* 18.2, pp. 283–328.
- Fu, Sascha, Mikko S Pakkanen, and Rama Cont (2024). “JAX-LOB: A GPU-accelerated limit order book simulator to unlock large scale reinforcement learning”. In: *arXiv preprint arXiv:2408.08806*.
- Garcez, Artur d’Avila and Luís C Lamb (2023). “Neurosymbolic AI: The 3rd wave”. In: *Artificial Intelligence Review*. DOI: [10.1007/s10462-023-10448-w](https://doi.org/10.1007/s10462-023-10448-w).
- George, Dileep et al. (2025). “A detailed theory of thalamic and cortical microcircuits for predictive visual inference”. In: *Science Advances*. DOI: [10.1126/sciadv.adr6698](https://doi.org/10.1126/sciadv.adr6698).
- Glasmachers, Tobias (2017). “Limits of end-to-end learning”. In: *PMLR*. Vol. 77.

- Goldblum, Micah et al. (2021). “Adversarial attacks on machine learning systems for high-frequency trading”. In: *ACM International Conference on AI in Finance (ICAIF)*. arXiv: 2002.09565.
- Halassa, Michael M and Sabine Kastner (2017). “Thalamic functions in distributed cognitive control”. In: *Nature Neuroscience* 20.12, pp. 1669–1679.
- Hassabis, Demis, Dharshan Kumaran, Christopher Summerfield, and Matthew Botvinick (2017). “Neuroscience-inspired artificial intelligence”. In: *Neuron* 95.2, pp. 245–258. DOI: [10.1016/j.neuron.2017.06.011](https://doi.org/10.1016/j.neuron.2017.06.011).
- Hugging Face (2024). *Candle: Minimalist ML framework for Rust*. URL: <https://github.com/huggingface/candle>.
- Internal Revenue Service (2024). *Wash sale rule basics for active traders*. URL: <https://www.irs.gov/taxtopics/tc409>.
- Jaskir, Adam and Michael J. Frank (2023). “On the normative advantages of dopamine and striatal opponency for learning and choice”. In: *eLife* 12, e85107.
- Joo, Hannah R. and Loren M. Frank (2023). “Selection of experience for memory by hippocampal sharp wave ripples”. In: *Science* 380. Sharp wave ripple selection mechanisms in hippocampus, pp. 1171–1175.
- Jung, Ralf, Jacques-Henri Jourdan, Robbert Krebbers, and Derek Dreyer (2018). “RustBelt: Securing the foundations of the Rust programming language”. In: *Proceedings of the ACM on Programming Languages* 2.POPL, p. 66. DOI: [10.1145/3158154](https://doi.org/10.1145/3158154).
- Kahneman, Daniel (2011). *Thinking, fast and slow*. Farrar, Straus and Giroux.
- Khandani, Amir E. and Andrew W. Lo (2011). “What happened to the quants in August 2007?” In: *Journal of Financial Markets* 14.1, pp. 1–46. DOI: [10.1016/j.finmar.2010.07.005](https://doi.org/10.1016/j.finmar.2010.07.005).
- Kralingen, Marc van, Diego Garlaschelli, Karolina Scholtus, and Iman van Lelyveld (2021). “Crowded trades, market clustering, and price instability”. In: *Entropy* 23.3, p. 336. DOI: [10.3390/e23030336](https://doi.org/10.3390/e23030336).
- Kumaran, Dharshan, Demis Hassabis, and James L. McClelland (2016). “What learning systems do intelligent agents need? Complementary learning systems theory updated”. In: *Trends in Cognitive Sciences* 20.7, pp. 512–534. DOI: [10.1016/j.tics.2016.05.004](https://doi.org/10.1016/j.tics.2016.05.004).
- LeCun, Yann, Yoshua Bengio, and Geoffrey Hinton (2015). “Deep learning”. In: *Nature* 521.7553, pp. 436–444.
- LeDoux, Joseph E (2000). “Emotion circuits in the brain”. In: *Annual Review of Neuroscience* 23, pp. 155–184. DOI: [10.1146/annurev.neuro.23.1.155](https://doi.org/10.1146/annurev.neuro.23.1.155).
- Lo, Andrew W. (2004). “The adaptive markets hypothesis”. In: *Journal of Portfolio Management* 30.5, pp. 15–29. DOI: [10.3905/jpm.2004.442611](https://doi.org/10.3905/jpm.2004.442611).

- Lo, Andrew W. (2017). *Adaptive Markets: Financial Evolution at the Speed of Thought*. Princeton University Press.
- López de Prado, Marcos (2018a). *Advances in Financial Machine Learning*. Wiley.
- (2018b). “The 10 reasons most machine learning funds fail”. In: *Journal of Portfolio Management* 44.6, pp. 120–133.
- Lou, Dong and Christopher Polk (2022). “Comomentum: Inferring arbitrage activity from return correlations”. In: *Review of Financial Studies* 35.7, pp. 3272–3302. DOI: [10.1093/rfs/hhab117](https://doi.org/10.1093/rfs/hhab117).
- Marino, Kevin et al. (2023). “ME-ViT: A single-load memory-efficient FPGA accelerator for vision transformers”. In: *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*.
- Markwick, David (2023). *Solving the Almgren-Chriss model*. URL: <https://www.actuaries.digital/2023/01/19/solving-the-almgren-chriss-model/>.
- Marra, Giuseppe, Sebastijan Dumančić, Robin Manhaeve, and Luc De Raedt (2024). “From statistical relational to neurosymbolic artificial intelligence: A survey”. In: *Artificial Intelligence* 328, p. 104062. DOI: [10.1016/j.artint.2023.104062](https://doi.org/10.1016/j.artint.2023.104062).
- Masset, Paul et al. (2025). “Multi-timescale reinforcement learning in the brain”. In: *Nature* 642, pp. 682–690. DOI: [10.1038/s41586-025-08929-9](https://doi.org/10.1038/s41586-025-08929-9).
- Matsakis, Nicholas D. and Il Klock Felix S. (2014). “The Rust language”. In: *ACM SIGAda Ada Letters* 34.3, pp. 103–104. DOI: [10.1145/2692956.2663188](https://doi.org/10.1145/2692956.2663188).
- Mattar, Marcelo G. and Nathaniel D. Daw (2018). “Prioritized memory access explains planning and hippocampal replay”. In: *Nature Neuroscience* 21, pp. 1609–1617.
- McClelland, James L, Bruce L McNaughton, and Randall C O’Reilly (1995). “Why there are complementary learning systems in the hippocampus and neocortex: insights from the successes and failures of connectionist models of learning and memory”. In: *Psychological Review* 102.3, p. 419. DOI: [10.1037/0033-295X.102.3.419](https://doi.org/10.1037/0033-295X.102.3.419).
- McInnes, Leland, John Healy, and James Melville (2018a). *UMAP: Uniform manifold approximation and projection (Python package)*. Includes AlignedUMAP for temporal manifold alignment. URL: <https://umap-learn.readthedocs.io/>.
- (2018b). “UMAP: Uniform manifold approximation and projection for dimension reduction”. In: *arXiv preprint arXiv:1802.03426*.
- Mnih, Volodymyr, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, et al. (2015). “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540, pp. 529–533. DOI: [10.1038/nature14236](https://doi.org/10.1038/nature14236).
- Mohan, A. et al. (2025). “Spiking neural network for cross-market portfolio optimization”. In: *arXiv preprint arXiv:2510.15921*.

- Monfils, Marie-H, Kiriana K Cowansage, Eric Klann, and Joseph E LeDoux (2009). “Extinction-reconsolidation boundaries: key to persistent attenuation of fear memories”. In: *Science* 324.5929, pp. 951–955.
- Mugel, Samuel, Enrique Lizaso, and Roman Orus (2022). “Dynamic portfolio optimization with real datasets using quantum processors and quantum-inspired tensor networks”. In: *Physical Review Research* 4, p. 013006. DOI: [10.1103/PhysRevResearch.4.013006](https://doi.org/10.1103/PhysRevResearch.4.013006).
- Padakandla, Sindhu et al. (2020). “Reinforcement learning algorithm for non-stationary environments”. In: *Applied Intelligence* 50. arXiv: 1905.03970, pp. 3590–3606.
- Polars Contributors (2024). *Polars: Lightning-fast DataFrame library*. URL: <https://www.pola.rs/>.
- Schaul, Tom, John Quan, Ioannis Antonoglou, and David Silver (2015). “Prioritized experience replay”. In: *arXiv preprint arXiv:1511.05952*.
- Schultz, Wolfram, Peter Dayan, and P. Read Montague (1997). “A neural substrate of prediction and reward”. In: *Science* 275.5306, pp. 1593–1599. DOI: [10.1126/science.275.5306.1593](https://doi.org/10.1126/science.275.5306.1593).
- Serafini, Luciano and Artur S. d’Avila Garcez (2016). “Logic tensor networks: Deep learning and logical reasoning from data and knowledge”. In: *AI*IA 2016*, pp. 334–348.
- Sokolov, Alexander A., R. Chris Miall, and Richard B. Ivry (2017). “The cerebellum: Adaptive prediction for movement and cognition”. In: *Trends in Cognitive Sciences* 21.5, pp. 313–332.
- Stein, Jeremy C. (2009). “Presidential address: Sophisticated investors and market efficiency”. In: *Journal of Finance* 64.4, pp. 1517–1548. DOI: [10.1111/j.1540-6261.2009.01472.x](https://doi.org/10.1111/j.1540-6261.2009.01472.x).
- Sterling, Peter (2012). “Allostasis: a model of predictive regulation”. In: *Physiology & Behavior* 106.1, pp. 5–15.
- Stillman, Namid R. and Rory Baggott (2024). “Neuro-symbolic traders: Assessing the wisdom of AI crowds in markets”. In: *arXiv preprint arXiv:2410.14587*. DOI: [10.48550/arXiv.2410.14587](https://doi.org/10.48550/arXiv.2410.14587).
- Sutton, Richard S. and Andrew G. Barto (2018). *Reinforcement Learning: An Introduction*. 2nd. MIT Press.
- Tokio Contributors (2024). *Tokio: Asynchronous runtime for Rust*. URL: <https://tokio.rs/>.
- Tóth, Bence, Yves Lemprière, Cyril Deremble, Joachim de Lataillade, Julien Kockelkoren, and Jean-Philippe Bouchaud (2011). “Anomalous price impact and the critical nature of liquidity in financial markets”. In: *Physical Review X* 1.2, p. 021006. DOI: [10.1103/PhysRevX.1.021006](https://doi.org/10.1103/PhysRevX.1.021006).

- Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin (2017). “Attention is all you need”. In: *Advances in Neural Information Processing Systems* 30.
- Vemeko (2023). *How to use FPGAs for HFT acceleration*. URL: <https://www.vemeko.com/fpga-hft>.
- Ven, Gido M. van de, Hava T. Siegelmann, and Andreas S. Tolias (2020). “Brain-inspired replay for continual learning with artificial neural networks”. In: *Nature Communications* 11, p. 4069.
- Volpati, Valerio, Michael Benzaquen, Zoltán Eisler, Iacopo Mastromatteo, Bence Tóth, and Jean-Philippe Bouchaud (2020). “Zooming in on equity factor crowding”. In: *arXiv preprint arXiv:2001.04185*.
- Wagner, Wolf (2011). “Systemic liquidation risk and the diversity–diversification trade-off”. In: *Journal of Finance* 66.4, pp. 1141–1175. DOI: [10.1111/j.1540-6261.2011.01666.x](https://doi.org/10.1111/j.1540-6261.2011.01666.x).
- Wan, Zishen et al. (2024). “Towards efficient neuro-symbolic AI: From workload characterization to hardware architecture”. In: *IEEE Transactions on Circuits and Systems for AI*. arXiv: 2409.13153.
- Wang, Zhiguang and Tim Oates (2015). “Imaging time-series to improve classification and imputation”. In: *arXiv preprint arXiv:1506.00327*. URL: <https://arxiv.org/abs/1506.00327>.
- Wilson, Hugh R and Jack D Cowan (1972). “Excitatory and inhibitory interactions in localized populations of model neurons”. In: *Biophysical Journal* 12.1, pp. 1–24.
- Wilson, Matthew A. and Bruce L. McNaughton (1994). “Reactivation of hippocampal ensemble memories during sleep”. In: *Science* 265.5172, pp. 676–679. DOI: [10.1126/science.8036517](https://doi.org/10.1126/science.8036517).
- Wolpert, Daniel M., R. Chris Miall, and Mitsuo Kawato (1998). “Internal models in the cerebellum”. In: *Trends in Cognitive Sciences* 2.9, pp. 338–347.
- Yamakawa, Hiroshi (2021). “The whole brain architecture approach: Accelerating the development of artificial general intelligence by referring to the brain”. In: *Neural Networks* 144, pp. 478–495.
- Zhang, Yiming et al. (2024). “Bidirectionally regulating gamma oscillations in Wilson-Cowan model by self-feedback loops: A computational study”. In: *Neural Computation*. Computational study of Wilson-Cowan oscillatory dynamics with bidirectional gamma regulation.